

Functional Pearl: A Tale of Two Provers

Verifying Parallel String Matching with Refinement and Dependent Types

ANONYMOUS AUTHOR(S)

Refinement Types have been extensively used to automatically verify shallow program properties, like safe indexing, but never before to prove deep, correctness theorems about realistic programs. We demonstrate for first time how refinement types are used for theorem proving by using Liquid Haskell, a refinement type checker for Haskell programs, to verify correctness of parallelization of a realistic, Haskell string matcher. We use refinement types to specify correctness properties, Haskell terms to express proofs of these properties, and Liquid Haskell to check the proofs. We evaluate Liquid Haskell as a theorem prover by replicating our 1428 LoC proof in the dependently-typed framework of the Coq proof assistant (1136 LoC); we compare both proofs, highlighting the relative advantages and disadvantages of the two provers.

1 INTRODUCTION

It was dependent types, it was refinement types.

Dependent type systems (like Coq [4] or Adga [17]) have been extensively used to specify theorems and proofs about programs and to machine-check the correctness of such proofs. These systems are equipped with a pool of theorems, tactics and proving environments that facilitate theorem proving. However, being geared towards verification, they provide minimal support for widespread features of general purpose languages like diverging computations, parallel environments or runtime-optimized libraries, some of which are only available via extraction.

Verification-oriented languages like Dafny [13], F* [25] and WhyML [9] combine good support for general purpose language features, like effectful and diverging computations, with semi-automated, SMT-based [2] verification. Focused on verification, these languages lack the non-verified but highly runtime optimized, real world libraries that come with existing general purpose languages. Moreover, all the above languages aim for highly expressive specifications, which makes SMT verification undecidable (in theory [16]) and unstable (in practice [14]).

Refinement Types [7, 11, 21] on the other hand, extend *existing* general purpose languages (including ML [3, 19, 30], C [6, 20], Haskell [28], Racket [12] and Scala [23]) with *predictable* SMT-based verification. Traditionally, to achieve predictable verification, refinement types were limited to shallow specifications. For example, we could use refinement types to specify that that appending two lists xs and ys yields a list of length equal to the sum of the lengths of xs and ys :

$$\text{append} :: xs:[a] \rightarrow ys:[a] \rightarrow \{v:[a] \mid \text{length } v = \text{length } xs + \text{length } ys\},$$

but not to express deeper properties, such as the associativity of `append`. This restriction critically limited the expressiveness of the specifications, but allowed for both automatic and predictable SMT-based [2] verification. Sadly, sadly, program equivalence proofs were beyond the expressive power of refinement types.

Liquid Haskell [28] extends refinement types with Refinement Reflection [29], a technique that reflects each function's implementation into the function's type, turning the refined language into a theorem prover while preserving predictable verification. In this paper we present the first non-trivial application of Liquid Haskell (1428 LoC) as a theorem prover, by proving the correctness of the parallelization of a naïve string matching

2017. This is the author's version of the work. It is posted here for your personal use. Not for redistribution. The definitive Version of Record was published in , <http://dx.doi.org/10.1145/nnnnnnnn.nnnnnnnn>.

algorithm based on an axiomatization of primitive parallel combinators. We replicate this proof in Coq (1136 LoC) and empirically compare the two approaches.

The contributions of this paper are:

- We explain how theorems and proofs are encoded and checked in Liquid Haskell by formalizing monoids and proving that lists form a monoid (§ 2), introducing notations and background necessary in the rest of the paper.
- We formalize monoid morphisms and show that such a morphism on a “chunkable” input can be correctly parallelized (§ 3) by:
 - (1) chunking up the input in chunks
 - (2) applying the morphism in parallel to all chunks, and
 - (3) recombining the mapped chunks using the monoid operation, also in parallel.
 We apply this result (§ 5) to obtain the first large application of Liquid Haskell as theorem prover: a verified parallelization of string matching.
- We evaluate the applicability of Liquid Haskell as a theorem prover by repeating the same proof in the Coq proof assistant. We identify interesting tradeoffs in the verification approaches encouraged by the two tools in two parts: we first draw preliminary conclusions based on the simpler parallelization theorem (§ 4) and then we delve deeper into the comparison, highlighting differences based on the string matching case study (§ 6). Finally, we complete the evaluation picture by providing additional quantitative comparisons (§ 7).

2 HASKELL FUNCTIONS AS PROOFS

Refinement Reflection [29] is a technique for writing Haskell functions which prove theorems about other Haskell functions and for machine-checking these proofs using Liquid Haskell [28]. In this section, as an introduction to Refinement Reflection, we prove that lists form a monoid by

- *specifying monoid laws* as refinement types,
- *proving the laws* by writing the implementation of the law specifications, and
- *verifying the proofs* using Liquid Haskell.

We start (§ 2.1) by defining a `List` datatype and the associated monoid elements ϵ and \diamond , corresponding to the empty list and concatenation. We then use Refinement Reflection [29] to prove the three monoid laws (§ 2.2, § 2.4, and § 2.5) in Liquid Haskell. To simplify the proofs, we use the tactic *PSE (Proof by Static Evaluation)* (§ 2.3) that automatically expands logic terms. Finally (§ 2.5), we conclude that lists are indeed monoids.

2.1 Reflection of Lists into Logic

To begin with, we define the standard recursive `List` datatype.

```
| data List [length] a = N | C {head :: a, tail :: List a}
```

The `length` annotation in the definition teaches Liquid Haskell to use the `length` function to check the termination of recursive list functions.

We define `length` as a standard Haskell function returning natural numbers.

```
| measure length
| length      :: List a → {v:Int | 0 ≤ v}
| length N    = 0
| length (C x xs) = 1 + length xs
```

But, what does Liquid Haskell know about `length`? Liquid Haskell enforces a clear separation between Haskell functions and their interpretation into the SMT logic allowing only the refinement specification of the function, *i.e.*, a decidable abstraction of the Haskell function, to flow into the SMT logic. With this separation Liquid Haskell achieves decidable and predictable type checking and prevents unstable program verification [14] encountered when arbitrary recursive functions flow into logic. However, this separation hinders precise verification. For instance, assume a list with 3 elements `xs = C 1 C 2 C 3 N`. Liquid Haskell can prove that the lengths of `xs` is a natural number, `xs :: {v : List a | 0 ≤ length xs}` but based on `length`'s specification alone, cannot prove that the length of `xs` is exactly 3 `xs :: {v : List a | 3 = length xs}`. To increase precision Liquid Haskell provides two mechanisms, **measure** and **reflect**, to carefully lift Haskell functions into logic, while preserving SMT-decidable program verification.

The **measure** annotation lifts `length` into the logic by strengthening the types of the `List` data constructors. For example, the type of `C` is strengthened to

$$| \text{C} : x : a \rightarrow xs : \text{List } a \rightarrow \{v : \text{List } a \mid \text{length } v = \text{length } xs + 1\}$$

where `length` is an uninterpreted function in the logic. In general, **measure** [28] annotations are used to precisely lift into the logic terminating, *unary* functions whose (1) domain is a data type and (2) body is a single case-expression over the datatype.

Then, we define and lift into logic the two monoid operators on Lists: an identity element ϵ (which is the empty list) and an associative operator \diamond (which is list append).

$ \begin{array}{l} \text{reflect } \diamond \\ (\diamond) :: \text{List } a \rightarrow \text{List } a \rightarrow \text{List } a \\ N \quad \quad \diamond \text{ } ys = ys \\ (C \ x \ xs) \diamond \text{ } ys = C \ x \ (xs \diamond \text{ } ys) \end{array} $	$ \begin{array}{l} \text{reflect } \epsilon \\ \epsilon :: \text{List } a \\ \epsilon = N \end{array} $
--	---

The **reflect** annotations lift \diamond and ϵ into logic by strengthening the types of the functions' specifications.

$ \begin{array}{l} (\epsilon) :: \{v : \text{List } a \mid v = \epsilon \wedge v = N\} \\ \\ (\diamond) :: xs : \text{List } a \rightarrow ys : \text{List } a \\ \quad \rightarrow \{v : \text{List } a \mid v = xs \diamond ys \wedge v = \text{if isN } xs \text{ then } ys \text{ else } C \ (\text{head } xs) \ (\text{tail } xs \diamond ys)\} \end{array} $	
---	--

where \diamond and ϵ are uninterpreted functions, and `isN`, `head` and `tail` are automatically generated measures. In general, **reflect** annotations are used to reflect terminating Haskell functions into the result of the function's type. After reflection, at each function call the function definition is unfolded exactly once into the logic, allowing Liquid Haskell to prove properties about Haskell functions.

2.2 Left Identity

In Liquid Haskell, we express theorems as refined type specifications and proofs as their Haskell inhabitants. We construct proof inhabitants using the combinators from the built-in `ProofCombinators` library that are summarized in Figure 1.

Left identity is expressed as a refinement type signature that takes as input a list `x : List a` and returns a **Proof** (*i.e.*, `unit`) type refined with the property $\epsilon \diamond x = x$.

$ \begin{array}{l} \text{idLeft_List} :: x : \text{List } a \rightarrow \{ \epsilon \diamond x = x \} \\ \text{idLeft_List } x = \epsilon \diamond x ==. N \diamond x ==. x \text{ *** QED} \end{array} $	
--	--

<pre> type Proof = () data QED = QED trivial :: Proof trivial = () (***) :: a -> QED -> Proof _ *** _ = () </pre>	<pre> (==.) :: x:a -> y:{a x = y} -> {v:a v = x} x ==. _ = x (∴) :: (Proof -> a) -> Proof -> a thm ∴ lemma = thm lemma (∧.) :: Proof -> Proof -> Proof _ ∧. _ = () </pre>
---	--

Fig. 1. Operators and Types defined in ProofCombinators. A **Proof** is a unit type that when refined is used to specify theorems. A trivial proof is the unit value. For example, `trivial :: {v:Proof | 1 + 2 = 3}` trivially proves the theorem $1 + 2 = 3$ by the SMT solver. `p *** QED` casts any expression `p` into a **Proof**. `x ==. y` asserts that `x` and `y` are equal. `thm ∴ lemma` proves `thm` using the lemma. `x ∧. y` combines two proofs `x` and `y` into one by inserting the argument proofs into the logical environment.

Here, $\{\epsilon \diamond x = x\}$ is a simplification for the **Proof** type $\{v:\mathbf{Proof} \mid \epsilon \diamond x = x\}$, since the binder `v` is irrelevant. We begin from the left hand side $\epsilon \diamond x$, which is equal to $N \diamond x$ by calling ϵ thus unfolding the equality `empty = N` into then logic. The proof combinator `x ==. y` let us equate `x` with `y` into the logic and returns `x` allowing us to continue the equational proof. Next, the call $N \diamond x$ unfolds into the logic the definition of (\diamond) on `N` and `x`, which is equal to `x`, concluding our proof. Finally, we use the operators `p *** QED` which casts `p` into a proof term. In short, the proof of left identity, proceeds by unfolding the definitions of ϵ and (\diamond) on the empty list.

2.3 PSE: Proof by Static Evaluation

PSE (Proof by Static Evaluation) is a terminating but incomplete heuristic (or tactic), inspired by [14], that Liquid Haskell uses to automatically unfold reflected functions in proof terms. PSE is not formalized in this paper but in § 7 we present how it can be used to simplify proof terms.

PSE can be used to simplify the left identity proof (we use the cornered one line code frame to denote Liquid Haskell proofs that use PSE).

```

idLeft_List :: x:List a → {  $\epsilon \diamond x = x$  }
idLeft_List _ = trivial

```

That is the proof proceeds, trivially, by symbolic evaluation of the expression $\epsilon \diamond x$.

2.4 Right Identity

Right identity is proved by structural induction. We encode inductive proofs by case splitting on the base and inductive case, and by enforcing the inductive hypothesis via a recursive call.

```

idRight_List :: x:List a → {  $x \diamond \epsilon = x$  }
idRight_List N = N  $\diamond$   $\epsilon$  ==. N *** QED

idRight_List (C x xs)
  = (C x xs)  $\diamond$   $\epsilon$ 
  ==. C x (xs  $\diamond$   $\epsilon$ )
  ==. C x xs ∴ idRight_List xs
  *** QED

```

The recursive call `idRight_List xs` is provided as a third optional argument in the `(==.)` operator to justify the equality $xs \diamond \epsilon = xs$, while the operator `(.·)` is merely a function application with the appropriate precedence. Note that LiquidHaskell, via termination and totality checking, is verifying that all the proof terms are well formed because (1) the inductive hypothesis is only applying to smaller terms and (2) all cases are covered.

We use the PSE tactic to automatically generate all function unfoldings and simplify the right identity proof.

```
[ idRight_List :: x:List a → { x ◇ ε = x }
  idRight_List N          = trivial
  idRight_List (C _ xs) = idRight_List xs
```

2.5 Associativity

Associativity is proved in a very similar manner, using structural induction.

```
[ assoc_List :: x:List a → y:List a → z:List a → { x ◇ (y ◇ z) = (x ◇ y) ◇ z }
  assoc_List N _ _          = trivial
  assoc_List (C _ x) y z = assoc_List xs y z
```

As with the left identity, the proof proceeds by (1) function unfolding (or rewriting in paper and pencil proof terms), (2) case splitting (or case analysis), and (3) recursion (or induction).

2.6 Lists are a Monoid

Finally, we formally define monoids as structures that satisfy the monoid laws of associativity and identity and conclude that `List a` is indeed a monoid.

Definition 2.1 (Monoid). The triple (m, ϵ, \diamond) is a monoid (with identity element ϵ and associative operator \diamond), if the following functions are defined.

```
idLeftm  :: x:m → {ε ◇ x = x}
idRightm :: x:m → {x ◇ ε = x}
assocm   :: x:m → y:m → z:m → {x ◇ (y ◇ z) = (x ◇ y) ◇ z}
```

THEOREM 2.2. $(List\ a, \epsilon, \diamond)$ is a monoid.

PROOF. `List a` is a monoid, as the implementation of `idLeft_List`, `idRight_List`, and `assoc_List` satisfy the specifications of the monoid laws `idLeftm`, `idRightm`, and `assocm`, with $m = List\ a$. \square

3 VERIFIED PARALLELIZATION OF MONOID MORPHISMS

A monoid morphism is a function between two monoids which preserves the monoidal structure, *i.e.*, preserves identity and distributes. We call a monoid morphism *chunkable* if its domain can be cut into chunks. A chunkable monoid morphism `f` is parallelized by:

- § 3.1 chunking up the input in chunks of size `i` (`chunk i`),
- § 3.2 applying the morphism in parallel to all chunks (`pmap f`), and
- § 3.3 recombining the chunks, in parallel `j` at a time, back to a single value (`pmconcat j`).

In this section we implement and verify in Liquid Haskell the correctness of the transformation

$$f = pmconcat\ j \cdot pmap\ f \cdot chunk\ i$$

This transformation relies on a single Haskell's library parallelization primitive that is *assumed* to be correct.

3.1 Lists are Chunkable Monoids

Definition 3.1 (Chunkable Monoids). We define a monoid (m, ϵ, \diamond) to be chunkable if for every natural number i and monoid x , the functions $\text{take}_m i x$ and $\text{drop}_m i x$ are defined in such a way as $\text{take}_m i x \diamond \text{drop}_m i x$ exactly reconstructs x .

```
lengthm :: m → Nat
dropm   :: i:Nat → x:{m | i ≤ lengthm x} → {v:m | lengthm v = lengthm x - i }
takem   :: i:Nat → x:{m | i ≤ lengthm x} → {v:m | lengthm v = i }

take_drop_specm :: i:Nat → x:m → {x = takem i x ⋄ dropm i x}
```

The functional methods of chunkable monoids are the take and drop while the length method is required to give the required pre- and post-condition on the other operations. Finally, take_drop_spec is a proof term that specifies the reconstruction property.

Next, we use the take_m and drop_m methods for each chunkable monoid (m, ϵ, \diamond) to define a $\text{chunk}_m i x$ function that splits x in chunks of size i .

```
chunkm :: i:Pos → x:m → {v:List m | chunk_specm i x v } / [lengthm x]
chunkm i x
  | lengthm x ≤ i = C x N
  | otherwise     = takem i x `C` chunkm i (dropm i x)
```

To prove termination of chunk_m Liquid Haskell checks that the user-defined termination metric $/ [\text{length}_m x]$ decreases at the recursive call. The check succeeds as $\text{drop}_m i x$ is specified to return a monoid smaller than x . We specify the length of the chunked result using the specification function chunk_spec_m .

```
chunk_specm i x v
  | lengthm x ≤ i = length v == 1
  | i == 1        = length v == lengthm x
  | otherwise     = length v < lengthm x
```

Liquid Haskell uses the specifications of both take_m and drop_m to automatically verify the length_m constraints imposed by chunk_spec_m .

Finally, we define lists from § ?? to be chunkable monoids according to the aforementioned specifications and we observe the benefits of SMT-based verification when reasoning about linear arithmetic.

<pre>take_List i N = N take_List i (C x xs) i == 0 = N otherwise = C x (take_List (i-1) xs)</pre>	<pre>drop_List i N = N drop_List i (C x xs) i == 0 = C x xs otherwise = drop_List (i-1) xs</pre>
--	---

The above definitions follow the library build-in definitions on lists, but, as noted earlier (§ ??), they need to be redefined for the reflected, user defined list data type. On the plus side, Liquid Haskell will *automatically* prove that the above definitions satisfy the specifications of the chunkable monoid on the length of § ?. Finally, the take-drop reconstruction specification is proved by induction on the size i and using the PSE tactic for the trivial static evaluation.

```

take_drop_spec_List i N
  = trivial
take_drop_spec_List i (C x xs) | i == 0
  = trivial
take_drop_spec_List i (C x xs)
  = take_drop_spec_List (i-1) xs

```

3.2 Parallel Map

We define a parallelized map function `pmap` using Haskell's library `parallel`. Concretely, we use the function `Control.Parallel.Strategies.withStrategy` that computes its argument in parallel given a parallel strategy.

```

pmap :: (a → b) → List a → List b
pmap f xs = withStrategy parStrategy (map f xs)

```

Parallelism in the Logic. The function `withStrategy`, that performs the runtime parallelization, is an imported Haskell library function, whose implementation is not available during verification. To use it in our verified code, we make the *assumption* that it always returns its second argument.

```

assume withStrategy :: Strategy a → x:a → {v:a | v = x}

```

Moreover, to reflect the implementation of `pmap` in the logic, `withStrategy` should also be represented in the logic. LiquidHaskell encodes `withStrategy` in the logic as a logical, *i.e.*, total, function that merely returns its second argument, `withStrategy _ x = x`. That is, our proof does not reason about parallelism in the logic. Rather, we assume correctness of the Haskell's library parallelization primitive.

Under this encoding, the strategy `parStrategy` does not affect verification. In our codebase we choose the traversable strategy.

```

parStrategy :: Strategy (List a)
parStrategy = parTraversable rseq

```

3.3 Parallel Monoidal Concatenation

The function `chunkm` lets us turn a monoidal value into several pieces. In the other direction, for any monoid (m, ϵ, \diamond) , the monoid concatenation `mconcatm` turns a chunked `List m` back into a single `m`.

```

mconcatm :: List m → m
mconcatm N      =  $\epsilon$ 
mconcatm (C x xs) = x  $\diamond$  mconcatm xs

```

Next, we parallelize the monoid concatenation by defining the function `pmconcatm` that chunks the input list of monoids and concatenates each chunk in parallel.

```

pmconcatm :: Int → List m → m
pmconcatm i x | i ≤ 1 || length x ≤ i
  = mconcatm x
pmconcatm i x
  = pmconcatm i (pmap mconcatm (chunk i x))

```

Where `chunk` is the list chunkable operation `chunk_List`. The function `pmconcatm i x` calls `mconcatm x` in the base case, otherwise it (1) chunks the list `x` in lists of size `i`, (2) runs in parallel `mconcatm` to each chunk, (3) recursively runs itself with the resulting list. Termination of `pmconcatm` holds, as the length of `chunk i x` is smaller than the length of `x`, when $1 < i$.

Finally, we prove correctness of parallelization of the monoid concatenation.

THEOREM 3.2. *For each monoid (m, ϵ, \diamond) the parallel and sequential concatenations are equivalent:*

```
| pmconcatEquivalence :: i:Int → x:List m → { pmconcatm i x = mconcatm x }
```

PROOF. We prove the theorem by providing an implementation of `pmconcatEquivalence` that satisfies its type. The proof proceeds by induction on `x`. The details of the proof can be found in [1], here we provide the sketch of the proof.

First, we prove that `mconcat` distributes over list splitting

```
| type LLEq m I = {List m | I ≤ length xs}
```

```
| mcut :: i:Nat → x:LLEq m i → {mconcatm x = mconcatm (take i x) ◇ mconcatm (drop i x)}
```

The proofs proceeds by structural induction, using monoid left identity in the base case and monoid associativity associativity and unfolding of `take` and `drop` methods in the inductive step.

We generalize the above lemma to prove that `mconcat` distributes over list chunking.

```
| mchunk :: i:Int → x:List m → {mconcatm x = mconcatm (map mconcatm (chunk i x))}
```

The proofs proceeds by structural induction, using monoid left identity in the base case and lemma `mconcatSplit` in the inductive step.

Lemma `mconcatChunk` is sufficient to prove `pmconcatEquivalence` by structural induction, using monoid left identity in the base case. □

3.4 Parallel Monoid Morphism

For any chunkable monoid `n`, monoid morphism `f :: n → m`, and natural number `i > 0` we can write a chunked version of `f` as

```
mconcat . pmap f . chunkn i :: n → m.
```

Before parallelizing `mconcat`, we will prove that the previous function is equivalent to `f`.

THEOREM 3.3 (MORPHISM DISTRIBUTION). *Let (m, ϵ, \diamond) be a monoid and (n, η, \square) be a chunkable monoid. Then, for every morphism $f : n \rightarrow m$, every positive number i and input x , $f x = mconcat (pmap f (chunk_n i x))$ holds.*

PROOF. We prove the theorem in Liquid Haskell we encode it into the specification of the function `morphism_distribution`.

```
| morphism_distribution :: f:(n → m) → Morphism n m f → x:n → i:Pos
    → {f x = mconcat (pmap f (chunkn i x))}
```

Where `Morphism` is a refinement type alias that captures the morphism properties from ??

```
| type Morphism n m F = x:n → y:n → {F η = ε ∧ F (x □ y) = F x ◇ F y}
```

We prove the theorem by providing an implementation of `morphism_distribution` that satisfies its type. The proof proceeds by induction on the length of the input.


```

morphismDistribution f thm x i
| lengthn x ≤ i
= idRightm (f is)
morphismDistribution f thm x i
= morphismDistribution f thm dropX i
  ∧. thm takeX dropX ∧. takeDropPropn i x
  where
    dropX = dropn i x
    takeX = taken i x

```

In the base case we use rewriting and right identity on the monoid $f \ x$. In the inductive case, we use the inductive hypothesis on the input $\text{dropX} = \text{drop}_n \ i \ x$, that is provably smaller than x as $1 < i$. Then, the fact that f is a monoid morphism, as encoded by our assumption argument thm takeX dropX we get basic distribution of f , that is $f \ \text{takeX} \diamond f \ \text{dropX} = f \ (\text{takeX} \sqcup \text{dropX})$. Finally, we merge $\text{takeX} \sqcup \text{dropX}$ to x using the property takeDropProp_n of the chunkable monoid n . \square

We can now replace the mconcat in our chunked monoid morphism in § ?? with pmconcat from § 3.3 to provide an implementation that uses parallelism to both map the monoid morphism and concatenate the results.

THEOREM 3.4 (CORRECTNESS OF PARALLELIZATION). *Let (m, ϵ, \diamond) be a monoid and (n, η, \sqcup) be a chunkable monoid. Then, for every morphism $f :: n \rightarrow m$, every positive numbers i and j , and input x , $f \ x = \text{pmconcat } i \ (\text{pmap } f \ (\text{chunk}_n \ j \ x))$ holds.*

PROOF.

We prove the theorem by providing an appropriate specification and implementation of the following `parallelismEquivalence` function.

```

parallelismEquivalence
:: f:(n → m) → Morphism n m f → x:n → i:Pos → j:Pos
→ {f x = pmconcat i (pmap f (chunkn j x))}

parallelismEquivalence f thm x i j
= pmconcatEquivalence i (pmap f (chunkn j x))
  ∧. morphismDistribution f thm x j

```

The proof follows merely by application of the two previous Theorems 3.3 and 3.2. \square

4 MONOID MORPHISM PARALLELIZATION IN COQ

To put using Liquid Haskell as a theorem prover in perspective, we replicated the same proofs using the Coq proof assistant. In this section we focus on the differences that appeared during proving the parallelization of monoid morphisms correct.

4.1 Intrinsic vs Extrinsic Verification

To begin with, we translate the chunkable monoid specification in Coq. The main difference is that the pre- and post-conditions of `take` and `drop` functions are not embedded in their types, but are independently encoded as specification terms in the extra `drop_spec` and `take_spec` methods.

```

lengthm : M → nat;
dropm   : nat → M → M;
takem   : nat → M → M;

drop_specm : ∀ i x, i ≤ lengthm x → lengthm (dropm i x) = lengthm x - i;
take_specm : ∀ i x, i ≤ lengthm x → lengthm (takem i x) = i;
take_drop_specm : ∀ i x, x = takem i x ∘ dropm i x;

```

The specification of chunkable monoids is a characteristic example of how Liquid Haskell and Coq naturally favor intrinsic and extrinsic verification, respectively. Liquid Haskell requires an intrinsic verification approach, as the (shallow) specifications of take and drop are embedded into the functions. On the contrary, Coq users can (and usually) take an extrinsic verification approach, where the specifications of take and drop are encoded as independent specification terms. We focus significantly more on the intrinsic/extrinsic distinction in Section 6, where we have more examples at hand.

4.2 Reasoning About Arithmetic

In Coq, we can leverage existing library functions on lists (here `ssreflect`'s `seq`) to define the chunkable monoid operations.

```

Definition length_list := @seq.size A;
Definition drop_list   := @seq.drop A;
Definition take_list   := @seq.take A;

```

We can also use library lemmata to prove most specifications of chunkable monoids directly.

```

Theorem drop_spec_list :
  ∀ i x, i ≤ length_list x → length_list (drop_list i x) = length_list x - i.
Proof. by apply seq.size_drop. Qed.

Theorem take_spec_list :
  ∀ i x, i ≤ length_list x → length_list (take_list i x) = i.
Proof.
  move ⇒ i x HSize.
  rewrite seq.size_take.
  destruct (i < size x) eqn:Size; ssromega.
Qed.

Theorem take_drop_spec_list :
  ∀ i x, x = takem i x ∘ dropm i x.
Proof.
  move ⇒ i x; symmetry; by apply seq.cat_take_drop.
Qed.

```

In all the above theorems, the crux of the proof lies in a library lemma (*i.e.*, `seq.size_take` for `take_spec_list`). The only case where the library theorems were not directly applicable reveals another interesting point of comparison between Liquid Haskell and Coq: dealing with arithmetic.

After branching on whether i is less than the size of x , we are left to prove that if i is less than or equal to x , but not strictly less than x , then the two numbers are necessarily equal.

```

A : Type
i : nat
x : seq A
HSize : i ≤ size x
Size : (i < size x) = false
=====
size x = i

```

Such an arithmetic fact would be discharged automatically in Liquid Haskell by the underlying SMT solver; in Coq, we must provide a proof term. Of course, this particular instance is very easy and Coq’s automation support can handle it given a sufficiently large hint database. However, the proof of the parallelization of monoid morphism relies heavily on a plethora of medium-sized linear arithmetic lemmas, a problem only exacerbated in the context of the string matching case study.

Most of the time we resort to an advanced Pressburger Arithmetic solver, omega, and in particular an adoption of it for `ssreflect`. However, when this automation doesn’t work (a situation which we encountered more than a couple of times), the prover must attempt to understand why it doesn’t and complete the proof. Worse, the proofs generated by the decision procedure are far from ideal. In the words of the Coq Reference Manual : “The simplification procedure is very dumb and this results in many redundant cases to explore. Much too slow.”

4.3 Non Structural Recursion

Since non-terminating programs would introduce a logical inconsistency, all Gallina programs are terminating. This is guaranteed through a particularly restrictive syntactic criterion: only recursive calls on subterms of some principal argument are allowed. To encode a computation that does not fall into this pattern, one solution would be to transform it to an equivalent one that does; we take such an approach in Section 6. However, such a translation is not always possible. Fortunately, the Coq community has invented techniques to get around this restriction. In particular, in Adam Chlipala’s CPDT, an entire chapter is devoted to three such general techniques, with complementary strengths and weaknesses: well-founded recursion, domain-theory-inspired non-termination monads and co-inductive non-termination monads.

For our proof effort we adopt the second approach: to keep the definition of `chunk` uncluttered from proof terms while persuading Coq that `chunk` terminates, we extended `chunk` with an additional natural number `fuel` argument that trivially decreases at each recursive call.

```

Fixpoint chunkm {M: Type} (fuel : nat) (i : nat) (x : M) : option (list M) :=
  match fuel with
  | 0 ⇒ None
  | S fuel' ⇒
    if lengthm x ≤ i then Some (cons x nil)
    else match chunkm fuel' i (dropm i x) with
      | Some res ⇒ Some (cons (takem i x) res)
      | None ⇒ None
    end
  end.

```

Thus `chunk` is defined to be `None` when not enough fuel is provided, otherwise it follows the above Haskell recursive definition. The fuel technique is common in Coq non-structural recursive definitions, heavily used for example in CompCert [?].

Our specification for `chunk` enforces both the same specification for the length of the result as Liquid Haskell (`chunk_resm`), and the (successful) termination of the computation given sufficient fuel:

```

Theorem chunk_specm : ∀ {M} i (x : M) ,
  i > 0 → exists l, chunkm (lengthm x).+1 i x = Some l /\ chunk_resm i x l.

```

A similar problem is encountered for `pmconcat` and addressed the same way.

4.4 General Purpose vs Verification Specific Languages

In Liquid Haskell, we can reason about actual Haskell programs using libraries from the Haskell ecosystem, like the one used for parallelization in the previous section. Coq does not have such a library, so instead of just axiomatizing the behavior of the library functions, we axiomatize the existence of such functions altogether:

```

Axiom Strategy      : Type.
Axiom parStrategy   : Strategy.
Axiom withStrategy   : ∀ {A}, Strategy → A → A.
Axiom withStrategy_spec : ∀ {A} (s : Strategy) (x : A), withStrategy s x = x.

```

In principle, we could extract these constants to their corresponding Haskell counterparts and obtain the same behavior.

5 CASE STUDY: CORRECTNESS OF PARALLEL STRING MATCHING IN LIQUID Haskell

In § 3 we presented that any monoid morphism whose domain is chunkable can be parallelized. We now make use of that result to parallelize string matching. We start by observing that strings are a chunkable monoid. We then turn string matching for a given target into a monoid morphism from a string to a suitable monoid, `SM target`, defined in § 5.2. Finally, in § 5.4, we parallelize string matching by a simple use of the parallel morphism function of § 3.4.

5.1 Refined Strings are assumed to be Chunkable Monoids

We define a new type `RString`, which is a chunkable monoid, to be the domain of our string matching function. Our type simply wraps Haskell's existing `ByteString`.

```

data RString = RS BS.ByteString

```

Similarly, we wrap the existing `ByteString` functions we will need to show `RString` is a chunkable monoid.

```

η = RS (BS.empty)
(RS x) ⊞ (RS y) = RS (x `BS.append` y)

lenStr    (RS x) = BS.length x
takeStr i (RS x) = RS (BS.take i x)
dropStr i (RS x) = RS (BS.take i x)

```

Although it is possible to explicitly prove that `ByteString` implements a chunkable monoid [27], it is time consuming and orthogonal to our purpose. Instead, we just *assume* the chunkable monoid properties of `RString`—thus demonstrating that refinement reflection is capable of doing gradual verification.

For instance, we define a logical uninterpreted function \sqsubseteq and relate it to the Haskell \sqsubseteq function via an assumed (unchecked) type.

```
| assume ( $\sqsubseteq$ ) :: x:RString  $\rightarrow$  y:RString  $\rightarrow$  {v:RString | v = x  $\sqsubseteq$  y}
```

Then, we use the uninterpreted function \sqsubseteq in the logic to assume monoid laws, like associativity.

```
| assume assocStr :: x:RString  $\rightarrow$  y:RString  $\rightarrow$  z:RString  $\rightarrow$  {x  $\sqsubseteq$  (y  $\sqsubseteq$  z) = (x  $\sqsubseteq$  y)  $\sqsubseteq$  z}
| assocStr _ _ = trivial
```

Haskell applications of \sqsubseteq are interpreted in the logic via the logical \sqsubseteq that satisfies associativity via theorem `assocStr`.

Similarly for the chunkable methods, we define the uninterpreted functions `takeStr`, `dropStr` and `lenStr` in the logic, and use them to strengthen the result types of the respective functions. With the above function definitions (in both Haskell and logic) and assumed type specifications, Liquid Haskell will check (or rather assume) that the specifications of chunkable monoid, as defined in the Definitions 2.1 and 3.1, are satisfied. We conclude with the assumption (rather than theorem) that `RString` is a chunkable monoid.

ASSUMPTION 5.1 (RSTRING IS A CHUNKABLE MONOID). *(RString, η , \sqsubseteq) combined with the methods `lenStr`, `takeStr`, `dropStr` and `takeDropPropStr` is a chunkable monoid.*

5.2 String Matching Monoid

String matching is determining all the indices in a source string where a given target string begins; for example, for source string `ababab` and target `aba` the results of string matching would be `[0, 2]`.

We now define a suitable monoid, `SM target`, for the codomain of a string matching function, where `target` is the string being looked for. Additionally, we will define a function `toSM :: RString \rightarrow SM target` which does the string matching and is indeed a monoid morphism from `RString` to `SM target` for a given `target`.

5.2.1 String Matching Monoid. We define the data type `SM target` to contain a refined string field `input` and a list of all the indices in `input` where the `target` appears.

```
| data SM (target :: Symbol) where
|   SM :: input:RString
|      $\rightarrow$  indices:[GoodIndex input target]
|      $\rightarrow$  SM target
```

We use the string type literal¹ to parameterize the monoid over the target being matched. This encoding allows the type checker to statically ensure that only searches for the same target can be merged together. The `input` field is a refined string, and the `indices` field is a list of good indices. For simplicity we present lists as Haskell's built-in lists, but our implementation uses the reflected list type, `L`, defined in § 2.

A `GoodIndex input target` is a refined type alias for a natural number `i` for which `target` appears at position `i` of `input`. As an example, the good indices of `"abcab"` on `"ababcbcab"` are `[2, 5]`.

```
| type GoodIndex Input Target
|   = {i:Nat | isGoodIndex Input (fromString Target) i }

| isGoodIndex :: RString  $\rightarrow$  RString  $\rightarrow$  Int  $\rightarrow$  Bool
| isGoodIndex input target i
|   = (subString i (lenStr target) input == target)
```

¹Symbol is a kind and target is effectively a singleton type.




Fig. 2. Mappend indices of String Matcher

```

    ^ (i + lenStr target ≤ lenStr input)

subString :: Int → Int → RString → RString
subString o l = takeStr l . dropStr o

```

5.2.2 Monoid Methods for String Matching. Next, we define the mappend and identity elements for string matching.

The *identity element* ϵ of $\text{SM } t$, for each target t , is defined to contain the identity $\text{RString } (\eta)$ and the identity $\text{List } ([])$.

```

 $\epsilon :: \forall (t :: \text{Symbol}). \text{SM } t$ 
 $\epsilon = \text{SM } \eta \ []$ 

```

The Haskell definition of \diamond , the monoid operation for $\text{SM } t$, is as follows.

```

( $\diamond$ ) ::  $\forall (t :: \text{Symbol}). \text{KnownSymbol } t \Rightarrow \text{SM } t \rightarrow \text{SM } t \rightarrow \text{SM } t$ 
( $\text{SM } x \text{ } xis$ )  $\diamond$  ( $\text{SM } y \text{ } yis$ )
=  $\text{SM } (x \sqcup y) (xis' ++ xyis ++ yis')$ 
  where
    tg  = fromString (symbolVal (Proxy :: Proxy t))
    xis' = map (castGoodIndexLeft tg x y) xis
    xyis = makeNewIndices x y tg
    yis' = map (shiftStringRight tg x y) yis

```

Note again that capturing target as a type parameter is critical, otherwise there is no way for the Haskell's type system to specify that both arguments of (\diamond) are string matchers on the same target.

The action of (\diamond) on the two input fields is straightforward; however, the action on the two indices is complicated by the need to shift indices and the possibility of new matches arising from the concatenation of the two input fields. Figure 2 illustrates the three pieces of the new indices field which we now explain in more detail.

1. Casting Good Indices. If xis is a list of good indices for the string x and the target tg , then xis is also a list of good indices for the string $x \sqcup y$ and the target tg , for each y . To prove this property we need to invoke the property `subStrAppendRight` on Refined Strings that establishes substring preservation on string right appending.

```

assume subStrAppendRight
  ::  $sl : \text{RString} \rightarrow sr : \text{RString} \rightarrow j : \text{Int}$ 
  →  $i : \{\text{Int} \mid i + j \leq \text{lenStr } sl\}$ 
  →  $\{ \text{subString } sl \ i \ j = \text{subString } (sl \sqcup sr) \ i \ j \}$ 

```

The specification of `subStrAppendRight` ensures that for each string `sl` and `sr` and each integer `i` and `j` whose sum is within `sl`, the substring from `i` with length `j` is identical in `sl` and in `(sl ⊔ sr)`. The function `castGoodIndexLeft` applies the above property to an index `i` to cast it from a good index on `sl` to a good index on `(sl ⊔ sr)`

```
castGoodIndexLeft
  :: tg:RString → sl:RString → sr:RString
  → i:GoodIndex sl tg
  → {v:GoodIndex (sl ⊔ sr) target | v = i}

castGoodIndexLeft tg sl sr i
  = cast (subStrAppendRight sl sr (lenStr tg) i) i
```

Where `cast p x` returns `x`, after enforcing the properties of `p` in the logic

```
cast :: b → x:a → {v:a | v = x }
cast _ x = x
```

Moreover, in the logic, each expression `cast p x` is reflected as `x`, thus allowing random (*i.e.*, non-reflected) Haskell expressions to appear in `p`.

2. Creation of new indices. The concatenation of two input strings `sl` and `sr` may create new good indices. For instance, concatenation of "ababcbab" with "cab" leads to a new occurrence of "abcbab" at index 5 which does not occur in either of the two input strings. These new good indices can appear only at the last `lenStr tg` positions of the left input `sl`. `makeNewIndices sl sr tg` detects all such good new indices.

```
makeNewIndices
  :: sl:RString → sr:RString → tg:RString
  → [GoodIndex {sl ⊔ sr} tg]
makeNewIndices sl sr tg
  | lenStr tg < 2 = []
  | otherwise    = makeIndices (sl ⊔ sr) tg lo hi
  where
    lo = maxInt (lenStr sl - (lenStr tg - 1)) 0
    hi = lenStr sl - 1
```

If the length of the `tg` is less than 2, then no new good indices are created. Otherwise, the call on `makeIndices` returns all the good indices of the input `sl ⊔ sr` for target `tg` in the range from `maxInt (lenStr sl - (lenStr tg - 1)) 0` to `lenStr sl - 1`.

Generally, `makeIndices s tg lo hi` returns the good indices of the input string `s` for target `tg` in the range from `lo` to `hi`.

```
makeIndices
  :: s:RString → tg:RString → lo:Nat
  → hi:Int → [GoodIndex s tg]
makeIndices s tg lo hi
  | hi < lo          = []
  | isGoodIndex s tg lo = lo:rest
  | otherwise        = rest
  where
```

```
rest = makeIndices s tg (lo + 1) hi
```

It is important to note that `makeNewIndices` does not scan all the input, instead only searching at most `lenStr tg` positions for new good indices. Thus, the time complexity to create the new indices is linear on the size of the target but independent of the size of the input.

3. *Shift Good Indices.* If `yis` is a list of good indices on the string `y` with target `tg`, then we need to shift each element of `yis` right `lenStr x` units to get a list of good indices for the string `x □ y`.

To prove this property we need to invoke the property `subStrAppendLeft` on Refined Strings that establishes substring shifting on string left appending.

```
assume subStrAppendLeft
  :: sl:RString → sr:RString
  → j:Int → i:Int
  → {subStr sr i j = subStr (sl □ sr) (lenStr sl+i) j}
```

The specification of `subStrAppendLeft` ensures that for each string `sl` and `sr` and each integers `i` and `j`, the substring from `i` with length `j` on `sr` is equal to the substring from `lenStr sl + i` with length `j` on `(sl □ sr)`. The function `shiftStringRight` both shifts the input index `i` by `lenStr sl` and applies the `subStrAppendLeft` property to it, casting `i` from a good index on `sr` to a good index on `(sl □ sr)`

Thus, `shiftStringRight` both appropriately shifts the index and casts the shifted index using the above theorem:

```
shiftStringRight
  :: tg:RString → sl:RString → sr:RString
  → i:GoodIndex sr tg
  → {v:(GoodIndex (sl □ sr) tg) | v = i + lenStr sl}
shiftStringRight tg sl sr i
  = subStrAppendLeft sl sr (lenStr tg) i
  'cast' i + lenStr sl
```

5.2.3 *String Matching is a Monoid.* Next we prove that the monoid methods ϵ and (\diamond) satisfy the monoid laws.

THEOREM 5.2 (SM IS A MONOID). *(SM t, ϵ , \diamond) is a monoid.*

PROOF. According to the Monoid Definition 2.1, we prove that string matching is a monoid, by providing safe implementations for the monoid law functions. First, we prove *left identity* using PSE.

```
idLeft :: x:SM t → { $\epsilon \diamond x = xs$  }
idLeft (SM i is)
  = idLeftStr i
  ∧. mapShiftZero tg i is
  ∧. newIsNullRight i tg
  ∧. idLeftList is
  where
    tg = fromString (symbolVal (Proxy :: Proxy t))
```

The proof proceeds by rewriting, using left identity of the monoid strings and lists, and two more lemmata.

- Identity of shifting by an empty string.


```

| mapShiftZero :: tg:RString → i:RString
  → is:[GoodIndex i target]
  → {map (shiftStringRight tg η i) is = is }

```

The lemma is proven by induction on *is* and the assumption that empty strings have length 0.

- No new indices are created.

```

| newIsNullLeft :: s:RString → t:RString
  → {makeNewIndices η s t = [] }

```

The proof relies on the fact that `makeIndices` is called on the empty range from 0 to -1 and returns [].

Next, we prove *right identity*.

```

| idRight :: x:SM t → {x ◊ ε = x }
idRight (SM i is)
  = idRightStr i
  ∧. mapCastId tg i η is
  ∧. newIsNullLeft i tg
  ∧. idRightList is
where
  tg = fromString (symbolVal (Proxy :: Proxy t))

```

The proof proceeds by rewriting, using right identity on strings and lists and two more lemmata.

- Identity of casting is proven

```

| mapCastId :: tg:RString → x:RString → y:RString
  → is:[GoodIndex x tg] →
  → {map (castGoodIndexRight tg x y) is = is}

```

We prove identity of casts by induction on *is* and identity of casting on a single index.

- No new indices are created.

```

| newIsNullLeft :: s:RString → t:RString
  → {makeNewIndices s η t = [] }

```

The proof proceeds by case splitting on the relative length of *s* and *t*. At each case we prove by induction that all the potential new indices would be out of bounds and thus no new good indices would be created.

- Finally we prove *associativity*. The PSE strategy failed to automatically prove associativity due to the complexity of the proof. For space, we only provide a proof sketch. The whole proof is available online [1]. Our goal is to show equality of the left and right associative string matchers.

```

| assoc :: x:SM t → y:SM t → z:SM t → { x ◊ (y ◊ z) = (x ◊ y) ◊ z }

```

To prove equality of the two string matchers we show that the input and indices fields are respectively equal. Equality of the input fields follows by associativity of RStrings. Equality of the index list proceeds in three steps.

- (1) Using list associativity and distribution of index shifting, we group the indices in the five lists shown in Figure 3: the indices of the input *x*, the new indices from mappending *x* to *y*, the indices of the input *y*, the new indices from mappending *x* to *y*, and the indices of the input *z*.




Fig. 3. Associativity of String Matching

- (2) The representation of each group depends on the order of appending. For example, if $zis1$ (resp. $zis2$) is the group zis when right (resp. left) mappend happened first, then we have

```
zis1 = map (shiftStringRight tg xi (yi ⊔ zi))
      (map (shiftStringRight tg yi zi) zis)

zis2 = map (shiftStringRight tg (xi ⊔ yi) zi) zis
```

That is, in right first, the indices of z are first shifted by the length of yi and then by the length of xi , while in the left first case, the indices of z are shifted by the length of $xi \sqcup yi$. In this second step of the proof we prove, using lemmata, the equivalence of the different group representations. The most interesting lemma we use is called `assocNewIndices` and proves equivalence of all the three middle groups together by case analysis on the relative lengths of the target tg and the middle string yi .

- (3) After proving equivalence of representations, we again use list associativity and distribution of casts to wrap the index groups back in string matchers.

We now sketch the three proof steps, while the whole proof is available online [1].

```
assoc x@(SM xi xis) y@(SM yi yis) z@(SM zi zis)
  -- Step 1: unwrapping the indices
  = x ⊔ (y ⊔ z)
  ==. (SM xi xis) ⊔ ((SM yi yis) ⊔ (SM zi zis))
  ...
  -- via list associativity and distribution of shifts
  ==. SM i (xis1 ++ ((xyis1 ++ yis1 ++ yzis1) ++ zis1))
  -- Step 2: Equivalence of representations
  ==. SM i (xis2 ++ ((xyis1 ++ yis1 ++ yzis1) ++ zis1))
    ∴ castConcat tg xi yi zi xis
  ==. SM i (xis2 ++ ((xyis1 ++ yis1 ++ yzis1) ++ zis2))
    ∴ mapLenFusion tg xi yi zi zis
  ==. SM i (xis2 ++ ((xyis2 ++ yis2 ++ yzis2) ++ zis2))
    ∴ assocNewIndices y tg xi yi zi yis
  -- Step 3: Wrapping the indices
  ...
  -- via list associativity and distribution of casts
  ==. (SM xi xis ⊔ SM yi yis) ⊔ SM zi zis
  = (x ⊔ y) ⊔ z
  *** QED
```

```

where
  i      = xi □ (yi □ zi)

  yzis1 = map (shiftStringRight tg xi (yi □ zi)) yzis
  yzis2 = makeNewIndices (xi □ yi) zi tg
  yzis  = makeNewIndices yi zi tg
  ...

```

□

5.3 String Matching Monoid Morphism

Next, we define the function $\text{toSM} :: \text{RString} \rightarrow \text{SM target}$ which does the actual string matching computation for a set target ²

```

toSM :: ∀ (target :: Symbol). (KnownSymbol target) ⇒ RString → SM target
toSM input = SM input (makeSMIndices input tg)
  where tg = fromString (symbolVal (Proxy :: Proxy target))

makeSMIndices :: x:RString → tg:RString → [GoodIndex x tg]
makeSMIndices x tg = makeIndices x tg 0 (lenStr tg - 1)

```

The input field of the result is the input string; the indices field is computed by calling `makeIndices` within the range of the input, that is from 0 to `lenStr input - 1`. We now prove that `toSM` is a monoid morphism.

THEOREM 5.3 (toSM IS A MORPHISM). $\text{toSM} :: \text{RString} \rightarrow \text{SM } t$ is a morphism between the monoids $(\text{RString}, \eta, \sqcup)$ and $(\text{SM } t, \epsilon, \diamond)$.

PROOF. Based on definition ??, proving `toSM` is a morphism requires constructing a valid inhabitant of the type

```

Morphism RString (SM t) toSM
= x:RString → y:RString → {toSM η = ε ∧ toSM (x □ y) = toSM x ◇ toSM y}

```

We define the function $\text{distributestoSM} :: \text{Morphism RString (SM } t) \text{ toSM}$ to be the required valid inhabitant.

The core of the proof starts from exploring the string matcher $\text{toSM } x \diamond \text{toSM } y$. This string matcher contains three sets of indices as illustrated in Figure 2: (1) x is from the input x , (2) xy is from appending the two strings, and (3) y is from the input y . We prove that appending these three groups of indices together gives exactly the good indices of $x \sqcup y$, which are also the value of the indices field in the result of $\text{toSM } (x \sqcup y)$.

```

distributestoSM x y
= (toSM x :: SM target) ◇ (toSM y :: SM target)
==. (SM x is1) ◇ (SM y is2)
==. SM i (xis ++ xyis ++ yis)
==. SM i (makeIndices i tg 0 hi1 ++ yis)
   ∴ (mapCastId tg x y is1 ∧ mergeNewIndices tg x y)
==. SM i (makeIndices i tg 0 hi1 ++ makeIndices i tg (hi1+1) hi)
   ∴ shiftIndicesRight 0 hi2 x y tg

```

²`toSM` assumes the target is clear from the calling context; it is also possible to write a wrapper function taking an explicit target which gets existentially reflected into the type.

```

==. SM i is
  ∴ mergeIndices i tg 0 hi1 hi
==. toSM (x ⊔ y)
*** QED
where
  xis = map (castGoodIndexRight tg x y) is1
  xyis = makeNewIndices x y tg
  yis = map (shiftStringRight tg x y) is2
  tg = fromString (symbolVal (Proxy::Proxy target))
  is1 = makeSMIndices x tg
  is2 = makeSMIndices y tg
  is = makeSMIndices i tg
  i = x ⊔ y
  hi1 = lenStr x - 1
  hi2 = lenStr y - 1
  hi = lenStr i - 1

```

The most interesting lemma we use is `mergeIndices x tg lo mid hi` that states that for the input `x` and the target `tg` if we append the indices in the range from `lo` to `mid` with the indices in the range from `mid+1` to `hi`, we get exactly the indices in the range from `lo` to `hi`. This property is formalized in the type of the lemma.

```

mergeIndices
  :: x:RString → tg:RString
  → lo:Nat → mid:{Int | lo ≤ mid} → hi:{Int | mid ≤ hi}
  → {makeIndices x tg lo hi = makeIndices x tg lo mid ++ makeIndices x tg (mid+1) hi}

```

The proof proceeds by induction on `mid` and using three more lemmata:

- `mergeNewIndices` states that appending the indices `xis` and `xyis` is equivalent to the good indices of `x ⊔ y` from `0` to `lenStr x - 1`. The proof case splits on the relative sizes of `tg` and `x` and is using `mergeIndices` on `mid = lenStr x1 - lenStr tg` in the case where `tg` is smaller than `x`.
- `mapCastId` states that casting a list of indices returns the same list.
- `shiftIndicesRight` states that shifting right `i` units the indices from `lo` to `hi` is equivalent to computing the indices from `i + lo` to `i + hi` on the string `x ⊔ y`, with `lenStr x = i`.

□

5.4 Parallel String Matching

We conclude this section with the definition of a parallelized version of string matching. We put all the theorems together to prove that the sequential and parallel versions always give the same result.

We define `toSMPar` as a parallel version of `toSM` using machinery of section 3.

```

toSMPar :: ∀ (target :: Symbol). (KnownSymbol target) ⇒ Int → Int → RString → SM
  target
toSMPar i j = pmconcat i . pmap toSM . chunkStr j

```

First, `chunkStr` splits the input into `j` chunks. Then, `pmap` applies `toSM` at each chunk in parallel. Finally, `pmconcat` concatenates the mapped chunks in parallel using \diamond , the monoidal operation for `SM target`.

Correctness. We prove correctness of `toSMPar` directly from Theorem 3.4.

THEOREM 5.4 (CORRECTNESS OF PARALLEL STRING MATCHING). *For each parameter i and j , and input x , `toSMPar i j x` is always equal to `toSM x`.*

```
| correctness :: i:Int → j:Int → x:RString → {toSM x = toSMPar i j x}
```

PROOF. The proof follows by direct application of Theorem 3.4 on the chunkable monoid $(RString, \eta, \sqcup)$ (by Assumption 5.1) and the monoid $(SM\ t, \epsilon, \diamond)$ (by Theorem 5.2).

```
| correctness i j x
|   = toSMPar i j x
|   ==. pmconcat i (pmap toSM (chunkStr j x))
|   ==. toSM is
|   ∴ parallelismEquivalence toSM distributestoSM x i j
|   *** QED
```

Note that application of the theorem `parallelismEquivalence` requires a proof that its first argument `toSM` is a morphism. By Theorem 3.3, the required proof is provided as the function `distributestoSM`. \square

6 STRING MATCHING IN COQ

In the previous section we saw in detail a proof of correctness for the parallelization of a string matching algorithm in Liquid Haskell. Just like in the monoid case, we replicated our proof in the Coq proof assistant. In this section we present the highlights of this effort, identifying more complementary strengths and weaknesses.

6.1 Existing Library Support

In the Liquid Haskell proof, we used a wrapper around `ByteStrings` to represent strings for efficiency; we also assumed the correctness of the `ByteString` operations instead of verifying them. In Coq, we used the built-in implementation of `Strings`. This allowed us to rely on the existing library theorems instead of assuming properties. We still admitted some theorems (e.g. the interoperation between `take` and `drop`), where direct counterparts were not available.

6.2 Proof inlining

In Liquid Haskell induction is encoded via recursive function calls. This highly restricts the structure of the Liquid Haskell proofs as each property proved by induction should be independently encoded as a recursive function/lemma. Thus, the user is forced to separately specify and proof each inductive lemma required for the proof. On the contrary, Coq does not impose any such restriction, allowing the user to prove lemmata inlined in the proof.

This convenience comes with the disadvantage that many times the proof repeatedly proves the same lemmata. For example, the `catIndices` property required to prove both distribution and associativity was expressed as a separate lemma in the Liquid Haskell proof, due to the lack of a different way to express it, while was by demand, inlined proven twice in the Coq proof.

6.3 Executable vs Inductive Specifications

In Liquid Haskell, `GoodIndex input tg` is a refinement type capturing the indices of the string input where the target string `tg` appears. Recall that `GoodIndex` is defined using the executable boolean predicate `isGoodIndex`.

```
| type GoodIndex Input Target = {i:Nat | isGoodIndex Input (fromString Target) i }
```

In fact, all refinements are executable boolean predicates written in Haskell. On the other hand, Coq users usually define *inductive* specifications which allow easier reasoning:

```

Definition isGoodIndex (input tg : string) (i : nat) :=
  (substring i (length tg) input) = tg /\ i + length tg ≤ length input.

```

However, in order to *test* whether a given index i is a good index for some given input and target strings, we need a decidability (*i.e.*, executable) procedure for `isGoodIndex`.

```

Definition isGoodIndexDec input tg i:
  {isGoodIndex input tg i} + {~ (isGoodIndex input tg i)}.
Proof.
  destruct (string_dec (substring i (length tg) input) tg) eqn:Eq;
  destruct (i + length tg ≤ length input) eqn:Ineq;
  auto; right ⇒ Contra; inversion Contra; eauto.
Qed.

```

Instead of returning a simple boolean, the decidability procedure returns a sum type:

```

Inductive sumbool (A B : Prop) : Set :=
  left : A → {A} + {B} | right : B → {A} + {B}

```

When extracted into OCaml or Haskell, `sumbool` is isomorphic to `Bool`; however, in Coq each constructor `left` and `right` carries additional proof information that can be used in proofs. This means that while the basic structure of the decidability procedure is straightforward (deciding whether both branches of the conjunction hold or not), it also contains additional content to construct appropriate proof terms.

6.4 Non Structural Recursion

Just like in § 5, string matching in Liquid Haskell requires an auxilliary recursive predicate `makeIndices` that uses non structural induction: a call to `makeIndices s tg lo hi` has a recursive call to `makeIndices s tg (lo+1) hi`, whose termination measure is $(hi - lo)$. In § 4 we dealt with non-structural recursion using `fuel`. In this case, it is easy to implement an equivalent predicate in Coq so that it *is* structurally recursive by calculating in advance how many steps (`cnt`) we need to take:

```

Fixpoint makeIndicesAux (s tg : string) (lo cnt : nat) : list nat :=
  match cnt with
  | 0 ⇒ nil
  | S cnt' ⇒
    let rest := makeIndicesAux s tg lo.+1 cnt' in
    if isGoodIndexDec s tg lo then lo::rest else rest
  end.

Definition makeIndices (s tg : string) (lo hi : nat) : list nat :=
  makeIndicesAux s tg lo (hi - lo).

```

6.5 Intrinsic vs Extrinsic Verification

There is an even more fundamental difference between the two `makeIndices` implementations: the type of the one in Coq does not mention any correctness properties unlike its Liquid Haskell counterpart! In order to

prove things in Liquid Haskell we had to use an intrinsic verification approach: by using refinement types the assumptions to our theorems are tied to the input types while the correctness results are tied to the output types. Thus, incorrect programs are impossible to even express. However, this approach has a couple of unfortunate drawbacks.

First, the computational content of the definitions must sometimes be cluttered to make the types work out. For example, when implementing the monoid operation \diamond for SM, we used `castGoodIndexLeft` to transform the refinement of the input (stating that every index in `xis` is a good index for `x`) to the refinement of the output (every index in `xis'` is a good index for `x++y`). However, computationally, this is identity: `xis'` is equal to `xis`. Second, *everything* we want to prove about a piece of code has to be proved at the same time. For instance, if we wanted to impose an additional constraint for good indices, we would have to revise all of our definitions, potentially adding more clutter.

In contrast, Coq users usually follow an extrinsic verification approach; operations over simply-typed expressions like `makeIndices` are proved correct after the fact:

```

Lemma makeIndicesAux_correct :
  ∀ cnt s tg lo,
    List.Forall (isGoodIndex s tg) (makeIndicesAux s tg lo cnt).
Proof.
  move ⇒ cnt; induction cnt ⇒ s tg lo // =;
  destruct (isGoodIndexDec s tg lo); simpl; auto.
Qed.

```

The use of an extrinsic approach in our proof development greatly simplifies the process. Specifically, the SM datatype is just a pair, while its validity is captured by a different inductive type.

```

Inductive SM (tg : string) :=
| Sm : ∀ (input : string) (l : list nat), SM tg.

Inductive validSM tg : SM tg → Prop :=
| ValidSM : ∀ input l,
  List.Forall (isGoodIndex input tg) l →
  validSM tg (Sm tg input l).

```

This extrinsic approach allows for cleaner implementations of the monoid operator \diamond_{sm} ,

```

Definition  $\diamond_{\text{sm}}$  tg (sm1 sm2 : SM tg) :=
  let '(Sm x xis) := sm1 in
  let '(Sm y yis) := sm2 in
  let xis' := xis in
  let xyis := makeNewIndices x y tg in
  let yis' := map (shiftStringRight tg x y) yis in
  Sm tg (x ++ y) (List.app xis' (List.app xyis yis')).

```

where `xis'` is *by definition* equal to `xis`. At the same time, the extrinsic approach clarifies exactly when the correctness assumptions for the index lists are necessary. For example, in the associativity proof of \diamond_{sm} we only require the middle string to be valid:

```

Theorem sm_assoc tg (sm1 sm2 sm3 : SM tg) : validSM tg sm2 →
   $\diamond_{\text{sm}}$  tg sm1 ( $\diamond_{\text{sm}}$  tg sm2 sm3) =  $\diamond_{\text{sm}}$  tg ( $\diamond_{\text{sm}}$  tg sm1 sm2) sm3.

```

On the other hand, the validity of \diamond_sm requires the validity of both inputs as preconditions:

```

Lemma sm_valid tg xs1 l1 xs2 l2 xs' l' :
  List.Forall (isGoodIndex xs1 tg) l1 →
  List.Forall (isGoodIndex xs2 tg) l2 →
   $\diamond\_sm$  tg (Sm tg xs1 l1) (Sm tg xs2 l2) = Sm tg xs' l' →
  List.Forall (isGoodIndex xs' tg) l'.

```

6.6 Dependent Pattern Matching

Unfortunately, we can not use extrinsic verification all the way through. A pair of strings and lists of natural numbers (like SM, the result of \diamond in Coq) is *not* a monoid by itself; only *valid* SMs form a monoid. To be able to use the monoid proofs of earlier sections we define a more restricted type `sm`, that carries along a proof of “goodness”.

```

Inductive sm tg : Type :=
| mk_sm :  $\forall$  xs l, List.Forall (isGoodIndex xs tg) l → sm tg.

```

Implementing the monoidal operation for this version of `sm` exemplifies the inconvenience of intrinsic approaches. Ideally, we would like to reuse the definition and properties of \diamond_sm directly, writing something like the following piece of code, where we would use `sm_valid` to construct the proof to fill `<proof>`.

```

 $\diamond$  sm1 sm2 :=
  match sm1, sm2 with
  | mk_sm xs1 l1 H1, mk_sm xs2 l2 H2 =>
    match  $\diamond\_sm$  tg (Sm tg xs1 l1) (Sm tg xs2 l2) with
    | Sm xs' l' => mk_sm tg xs' l' <proof>
  end
end

```

However, dependent pattern matching in Coq does not by default provide an equality between \diamond_sm tg (Sm tg xs1 l1) (Sm tg xs2 l2) and Sm xs' l' in scope for `<proof>`. Instead, the user must resort to what is known as the *convoy pattern*: the result of the inner match becomes a function that takes evidence of the needed equality as an argument, while the entire match is applied to such evidence.

```

 $\diamond$  sm1 sm2 :=
  match sm1, sm2 with
  | mk_sm xs1 l1 H1, mk_sm xs2 l2 H2 =>
    let s :=  $\diamond\_sm$  tg (Sm tg xs1 l1) (Sm tg xs2 l2) in
    let App := erefl s in
    (match s as s0 return ( $\diamond\_sm$  tg (Sm tg xs1 l1) (Sm tg xs2 l2) = s0 → sm tg)
    with
    | Sm xs' l' => fun _ => mk_sm tg xs' l' _
    end) App
  end
end

```

6.7 Proof Irrelevance

But we have an even bigger problem. Unlike Liquid Haskell where the intrinsic specifications only live at the logic level, in Coq they are part of the terms. Which means that the associativity proof of \diamond requires that all of

Property	Coq				Liquid Haskell				Liquid Haskell + PSE			
	Time	Spec	Proof	Exec	Time	Spec	Proof	Exec	Time	Spec	Proof	Exec
Parallelization	5	121	329	39	8	54	164	78	5	62	73	78
String Matcher	33	127	437	83	87	199	831	102	1287	223	596	102
Total	38	248	766	122	95	253	995	180	1292	285	669	180

Table 1. Quantitative evaluation of the proofs. We report verification time and LoC (Lines of Code) required to prove the general **parallelization** equivalence of monoid morphisms and its application to the **string matcher**. We split the proofs of Coq (1136 LoC in total), Liquid Haskell (1428 LoC in total) and Liquid Haskell with PSE (1134 LoC in total) into **specifications**, **proof terms** and **executable** code. **Time** is verification time in seconds.

the string, integer list and validity proof components of the resulting sms are *syntactically* equal. However, such proofs are *not* necessarily equal!

To that end, we use *Proof Irrelevance*, an admissible axiom, consistent with Coq’s logic (but not necessarily other axioms like *univalence*), which states that any two proofs of the same property are equal.

```
|| proof_irrelevance : ∀ (P : Prop) (p1 p2 : P), p1 = p2
```

For sanity check, we only use it once to prove an equality lemma for sms, that only require the string and integer list components to be equal.

```
|| Lemma proof_irrelevant_equality tg xs xs' l H l' H' :
||   xs = xs' → l = l' → mk_sm tg xs l H = mk_sm tg xs' l' H'.
```

Using the proof irrelevant equality we were able to prove the monoid instance of `sm` (and similar tricks were necessary for `monoidmorphism`). We conjecture that this proof is impossible without using such an axiom.

7 EVALUATION

7.1 Quantitative Comparison.

Table 1 summarizes the quantitative evaluation of our two proofs: the generalized equivalence property of parallelization of monoid morphisms and its application on the parallelization of a naïve string matcher. We used three provers to conduct our proofs: Coq, Liquid Haskell, and Liquid Haskell extended with the PSE (Proof by Static Evaluation § 2.3) heuristic. The Liquid Haskell proof was originally specified and verified by a Liquid Haskell expert within 2 months. Most of this time was spend on iterating between incorrect implementations of the string matching implementation (and the proof) based on Liquid Haskell’s type errors. After the Liquid Haskell proof was finalized, it was ported to Coq by a Coq expert within 2 weeks. We note that, due to the complexity of the proofs, none of them were optimized neither for size nor for verification time.

Verification time. We verified our proofs using a machine with an Intel Core i7-4712HQ CPU and 16GB of RAM. Verification in Coq is the fastest requiring 38 sec in total. Liquid Haskell requires x2.5 as much time while it needs x34 time using PSE. This slowdown is expected given that, unlike Coq that is checking the proof, Liquid Haskell uses the SMT solver to synthesize proof terms during verification, while PSE is an under-development, non-optimized approach to heuristically synthesize proof terms by static evaluation. In small proofs, like the generalized parallelization theorem, PSE can speedup verification time as proofs are quickly synthesized due to the fewer reflected functions and proof terms are much smaller leading to faster Liquid Haskell verification.

Verification size. We split the total numbers of code into three categories for both Coq and Liquid Haskell.

- **Spec** represents the theorem and lemma definitions, and the refinement type specifications, resp..

- **Proofs** represents the proof scripts and the Haskell proof terms (*i.e.*, **Proof** resulting functions), resp..
- **Exec** represents in both provers the executable portion of the code.

Counting both specifications and proofs as verification code, we conclude that in Coq the proof requires 8x the lines of the executable code. This ration drops to 7x for Liquid Haskell, because the executable code in the Haskell implementation is increased to include a basic string matching interface for printing and testing the application. Finally, the ration drops to 5x when the PSE heuristic is used, as the proof terms are shrunk without any modification to the executable portion.

Evaluation of PSE. PSE is used to synthesize non-sophisticated proof terms, leading to fewer lines of proof code but slower verification time. We used PSE to synthesize 31 out of the 43 total number of proof terms. PSE failed to synthesize the rest proof terms due to: 1. *incompleteness*: PSE is unable to synthesize proof terms when the proof structure does not follow the structure of the reflected functions, or 2. *verification slowdown*: in big proof terms there are many intermediate terms to be evaluated which dreadfully slows verification. Formalization and optimization of PSE, so that it synthesizes more proof terms faster, is left as future work.

7.2 Qualitative Comparison.

We summarize the essential differences in theorem proving using Liquid Haskell versus Coq based on our experience (§ 4 and § 6). These differences validate and illustrate the distinctions that have been previously [5, 19, 25] described between Refinement and Dependent Types.

General Purpose vs. Verification Specific Languages. Haskell is a general purpose language with highly tuned parallelism and optimized libraries (*e.g.*, `ByteString`) that we used (§ ??) to build a real applications, that moreover is provably correct. Coq lacks such features providing no direct support for parallel or divergence computations. But, Coq is a verification specific language, equipped with a pool of theorems (*e.g.*, list and string associativity) and tactics that ease theorem proving (§ 6.1). Liquid Haskell has no library theorems forcing us to prove all theorems from scratch. Finally, Coq was designed as a verification tool with a minimum trusted code base. On the contrary, Liquid Haskell can be used a theorem prover assuming one trusts GHC's type inference, Liquid Haskell's constrain generation, and SMT's validity checking.

SMT-automation vs. Tactics. Liquid Haskell uses an SMT-solver to automate proofs over decidable theories (such as linear arithmetic, uninterpreted functions), while Coq requires explicit proofs via tactics (§ ??). SMT-automation highly reduces the proof burden but increases verification time as proof synthesis occurs during type checking.

Intrinsic and Extrinsic verification. Liquid Haskell naturally uses intrinsic verification, *i.e.*, specifications are embedded in the definitions of the functions, should be proven (automatically by SMTs) at function definitions and are assumed at function calls. Coq uses extrinsic verification that separates the functionality of definitions from their specifications that can be independently proven and explicitly called (§ ??). Extrinsic verification makes function definitions cleaner in Coq, but when data depend on their specifications, as in the string matching case study, verification requires advanced features like Proof Irrelevance § 6.7 and Dependent Pattern Matching § 6.6.

Semantic vs. Structural Termination Checking. Liquid Haskell uses a semantics termination checker that proves termination given a wellfounded termination metric. On the contrary, Coq allows inductive functions to be defined only on an argument that is structurally decreasing. This constrain simplifies Coq's metatheory but litters non-trivially inductive functions with an extra fueling parameter (§ ??, § 6.4).

8 RELATED WORK

SMT-Based Verification. SMT solvers have been extensively used to automate reasoning on verification languages like Dafny [13], Fstar [25] and Why3 [9]. These languages are designed for verification, thus have limited support for commonly used language features like parallelism and optimized libraries that we use in our verified implementation. Refinement Types [7, 11, 21] on the other hand, target existing general purpose languages, such

as ML [3, 19, 30], C [6, 20], Haskell [28], Racket [12] and Scala [23]. However, before Refinement Reflection [29] was introduced, they only allowed “shallow” program specifications, that is, properties that only talk about behaviors of program functions but not functions themselves.

Dependent Types. Unlike Refinement Types, dependent type systems, like Coq [4], Adga [17] and Isabelle/HOL [18] allow for “deep” specifications which talk about program functions, such as the program equivalence reasoning we presented. Compared to (Liquid) Haskell, these systems allow for tactics and heuristics that automate proof term generation but lack SMT automations and general-purpose language features, like non-termination, exceptions and IO. Zombie [5, 24] and Fstar [25] allow dependent types to coexist with divergent and effectful programs, but still lack the optimized libraries, like `ByteString`, which come with a general purpose language with long history, like Haskell.

Parallel Code Verification. Dependent type theorem provers have been used before to verify parallel code. BSP-Why [10] is an extension to Why2 that is using both Coq and SMTs to discharge user specified verification conditions. Daum [8] used Isabelle to formalize the semantics of a type-safe subset of C, by extending Schirmer’s [22] formalization of sequential imperative languages. Finally, Swierstra [26] formalized mutable arrays in Agda and used them to reason about distributed maps and sums.

One work closely related to ours is SyDPaCC [15], a Coq library that automatically parallelizes list homomorphisms by extracting parallel Ocaml versions of user provided Coq functions. Unlike SyDPaCC, we are not automatically generating the parallel function version, because of engineering limitations (§ 7). Once these are addressed, code extraction can be naturally implemented by turning Theorem 3.4 into a Haskell type class with a default parallelization method. SyDPaCC used maximum prefix sum as a case study, whose morphism verification is much simpler than our string matching case study. Finally, our implementation consists of 2K lines of Liquid Haskell, which we consider verbose and aim to use tactics to simplify. On the contrary, the SyDPaCC implementation requires three different languages: 2K lines of Coq with tactics, 600 lines of Ocaml and 120 lines of C, and is considered “very concise”.

9 CONCLUSION

We described how Liquid Haskell equipped with Refinement Reflection can be used as a theorem prover by presenting its first non-trivial application to a realistic Haskell program: parallelization of a string matcher. We ported our 1428 LoC proof to the Coq proof assistant (1136 LoC) and compared the two provers capturing the essential differences of using dependent and refinement types for theorem proving. We conclude that the strong points of Liquid Haskell as a theorem prover is that the proof refers to executable Haskell code that directly uses advanced Haskell’s features like optimized libraries, parallel or diverging code and that the proof is SMT-automated over decidable theories (like linear arithmetic). The strong points of Coq is that the proof is checked assuming a minimum trusted code base and proof development is assisted by a pool of library theorems and tactics. [Make some reference to the book here.](#)

REFERENCES

- [1] Code for verified string indexing. 2017. Provided in Non-anonymous Supplementary Material.
- [2] C. Barrett, A. Stump, and C. Tinelli. The SMT-LIB Standard: Version 2.0. 2010.
- [3] J. Bengtson, K. Bhargavan, C. Fournet, A.D. Gordon, and S. Maffei. Refinement types for secure implementations. In *CSF*, 2008.
- [4] Y. Bertot and P. Castéran. *Coq’Art: The Calculus of Inductive Constructions*. Springer Verlag, 2004.
- [5] C. Casinghino, V. Sjöberg, and S. Weirich. Combining proofs and programs in a dependently typed language. In *POPL*, 2014.
- [6] Jeremy Condit, Matthew Harren, Zachary R. Anderson, David Gay, and George C. Necula. Dependent types for low-level programming. In *ESOP*, 2007.
- [7] R. L. Constable and S. F. Smith. Partial objects in constructive type theory. In *LICS*, 1987.
- [8] M Daum. Reasoning on Data-Parallel Programs in Isabelle/Hol. In *C/C++ Verification Workshop*, 2007.

- [9] Jean-Christophe Filliâtre and Andrei Paskevich. Why3 – Where Programs Meet Provers. In *ESOP*, 2013.
- [10] J Fortin and F. Gava. BSP-Why: A tool for deductive verification of BSP algorithms with subgroup synchronisation. In *Int J Parallel Prog.* 2015.
- [11] T. Freeman and F. Pfenning. Refinement types for ML. In *PLDI*, 1991.
- [12] Andrew M. Kent, David Kempe, and Sam Tobin-Hochstadt. Occurrence typing modulo theories. In *PLDI*, 2016.
- [13] K. Rustan M. Leino. Dafny: An automatic program verifier for functional correctness. *LPAR*, 2010.
- [14] Rustan Leino and Clément Pit-Claudel. Trigger selection strategies to stabilize program verifiers. 2016.
- [15] Frédéric Loulergue, Wadoud Bousdira, and Julien Tesson. Calculating Parallel Programs in Coq using List Homomorphisms. In *International Journal of Parallel Programming*, 2016.
- [16] Michał Moskal, Jakub Lopuszański, and Joseph R. Kiniry. E-matching for Fun and Profit. In *Electron. Notes Theor. Comput. Sci.*, 2008.
- [17] U. Norell. *Towards a practical programming language based on dependent type theory*. PhD thesis, Chalmers, 2007.
- [18] L. C. Paulson. Isabelle fi?! A Generic Theorem prover. *Lecture Notes in Computer Science*, 1994.
- [19] P. Rondon, M. Kawaguchi, and R. Jhala. Liquid types. In *PLDI*, 2008.
- [20] P. Rondon, M. Kawaguchi, and R. Jhala. Low-level liquid types. In *POPL*, 2010.
- [21] J. Rushby, S. Owre, and N. Shankar. Subtypes for specifications: Predicate subtyping in pvs. *IEEE TSE*, 1998.
- [22] N Schirmer. *Verification of Sequential Imperative Programs in Isabelle/HOL*. PhD thesis, TU Munich, 2006.
- [23] Georg Stefan Schmid and Viktor Kuncak. SMT-based Checking of Predicate-Qualified Types for Scala. In *Scala*, 2016.
- [24] Vilhelm Sjöberg and Stephanie Weirich. Programming up to congruence. *POPL*, 2015.
- [25] Nikhil Swamy, Cătălin Hrițcu, Chantal Keller, Aseem Rastogi, Antoine Delignat-Lavaud, Simon Forest, Karthikeyan Bhargavan, Cédric Fournet, Pierre-Yves Strub, Markulf Kohlweiss, Jean-Karim Zinzindohoue, and Santiago Zanella-Béguelin. Dependent types and multi-monadic effects in F*. In *POPL*, 2016.
- [26] Wouter Swierstra. More dependent types for distributed arrays. 2010.
- [27] N. Vazou, E. L. Seidel, and R. Jhala. Liquidhaskell: Experience with refinement types in the real world. In *Haskell Symposium*, 2014.
- [28] N. Vazou, E. L. Seidel, R. Jhala, D. Vytiniotis, and S. Peyton-Jones. Refinement Types for Haskell. In *ICFP*, 2014.
- [29] Niki Vazou and Ranjit Jhala. Refinement Reflection. arXiv:1610.04641, 2016.
- [30] H. Xi and F. Pfenning. Eliminating array bound checking through dependent types. In *PLDI*, 1998.