Towards Complete Specification and Verification with SMT

ANONYMOUS AUTHOR(S)

We introduce *Refinement Reflection*, a new framework for building SMT-based deductive verifiers. The key idea is to reflect the code implementing a user-defined function into the function's (output) refinement type. As a consequence, at *uses* of the function, the function definition is instantiated in a precise fashion that permits decidable verification. We show how reflection allows the user to write *equational proofs* of programs just by writing other programs *e.g.* using pattern-matching and recursion to perform case-splitting and induction. Thus, via, the Curry-Howard correspondence we show that reflection permits the *specification* of arbitrary functional correctness properties. While equational proofs are elegant, writing them out can be exhausting. We introduce a proof-search algorithm called *Proof by Logical Evaluation* that uses techniques from model checking & abstract interpretation, to completely automate equational reasoning. We have implemented reflection in LIQUID HASKELL and used it to verify that the widely used instances of the Monoid, Applicative, Functor, and Monad typeclasses actually satisfy key algebraic laws required to make the clients safe, and to build the first library that actually verifies assumptions about associativity and ordering that are crucial for safe deterministic parallelism.

1 INTRODUCTION

Deductive verifiers fall roughly into two camps. Satisfiability Modulo Theory (SMT) based verifiers (*e.g.* DAFNY and F*) use fast decision procedures to completely automate the verification of programs that only require reasoning over a fixed set of theories like linear arithmetic, string, set and bitvector operations. These verifiers, however, encode the semantics of user-defined functions with universally-quantified axioms and use incomplete (albeit effective) heuristics to instantiate those axioms. These heuristics make it difficult to characterize the kinds of proofs that can be automated, and hence, explain why a given query fails [Leino and Pit-Claudel 2016]. At the other extreme, we have Type-Theory (TT) based verifiers like CoQ and AGDA that use type-level computation (normalization) to facilitate principled reasoning about total, terminating user-defined functions. Unfortunately, these verifiers require the user to supply lemmas or rewrite hints that add friction to the ubiquitous task of discharging simple queries over decidable theories.

We introduce *Refinement Reflection*, a new framework for building SMT-based deductive verifiers, which permits the specification of arbitrary properties and yet enables complete, automated SMT-based reasoning about user-defined functions. In previous work, refinement types [Constable and Smith 1987; Rushby et al. 1998] — which decorate basic types (*e.g.* `Integer`) with SMT-decidable predicates (*e.g.* {v : `Integer` | 0 ≤ v && v < 100}) — were used to retrofit so-called shallow verification, such as array bounds checking, into several languages: ML [Bengtson et al. 2008; Rondon et al. 2008; Xi and Pfenning 1998], C [Condit et al. 2007; Rondon et al. 2010], Haskell [Vazou et al. 2014], TypeScript [Vekris et al. 2016], and Racket [Kent et al. 2016].

1. Refinement Reflection Our first contribution is the notion of *refinement reflection*. To reason about user-defined functions, the function's implementation can be *reflected* into its (output)

refinement-type specification, thus converting the function's type signature into a precise description of the function's behavior. This simple idea has a profound consequence: at *uses* of the function, the standard rule for (dependent) function application yields a precise means of reasoning about the function (§ 4).

2. Complete Specification Reflection lets us represent proofs simply as unit-values refined by the proposition that they prove. Our second contribution is a *library of combinators* that lets programmers compose sophisticated *equational proofs* from basic refinements and function definitions. Our proof combinators let programmers use existing language mechanisms like branches (to encode case splits), recursion (to encode induction), and functions (to encode auxiliary lemmas) to write proofs that look very much like transcriptions of their pencil-and-paper analogues (§ 2). Furthermore, since proofs are literally just programs, we demonstrate, via the Curry-Howard Isomorphism, that Refinement Reflection yields a proof system that is equivalent to (§ 3)

3. Complete Verification While equational proofs can be very elegant and expressive, writing them out can quickly get exhausting. Our third contribution is *Proof by Logical Evaluation* (PLE) a new proof-search algorithm that completely automates equational reasoning. The key idea in PLE is to mimic type-level computation within SMT-logics by representing functions in a *guarded form* [Dijkstra 1975] and repeatedly unfolding function application terms by instantiating them with their definition corresponding to an *enabled* guard. We formalize a notion of equational proof and show that the above strategy is *complete*: *i.e.* it is guaranteed to find an equational proof if one exists. Furthermore, using techniques from the literature on Abstract Interpretation [Cousot and Cousot 1977] and Model Checking [Clarke et al. 1992], we show that the above proof search corresponds to a *universal* (or *must*) abstraction of the concrete semantics of the user-defined functions. Thus, as those functions are total we obtain the pleasing guarantee that proof search terminates (§ 6).

We evaluate our approach by implementing refinement reflection and PLE in Liquid Haskell [Vazou et al. 2014], thereby turning Haskell into a theorem prover. Repurposing an existing programming language allows us to take advantage of a mature compiler and an ecosystem of libraries, while keeping proofs and programs in the same language. We demonstrate the benefits of this conversion by proving typeclass laws. Haskell's typeclass machinery has led to a suite of expressive abstractions and optimizations which, for correctness, crucially require typeclass *instances* to obey key algebraic laws. We show how reflection and PLE can be used to verify that widely used instances of the Monoid, Applicative, Functor, and Monad typeclasses satisfy the respective laws. Finally, we use reflection to create the first deterministic parallelism library that actually verifies assumptions about associativity and ordering that ensure determinism (§ 7).

Thus, our results demonstrate that Refinement Reflection and Proof by Logical Evaluation identify a new design for deductive verifiers which, by combining the complementary strengths of SMT- and TT- based approaches, enables complete verification of expressive specifications spanning decidable theories and user defined functions.

2 OVERVIEW

We start with an overview of how SMT-based refinement reflection lets us write equational proofs as plain functions and how PLE automates equational reasoning.

2.1 Refinement Types

First, we recall some preliminaries about specification and verification with refinement types.

Refinement types are the source program's (here Haskell's) types refined with logical predicates drawn from an SMT-decidable logic [Constable and Smith 1987; Rushby et al. 1998]. For example, we

define `Nat` as the set of `Integer` values v that satisfy the predicate $0 \leq v$ from the quantifier-free logic of linear arithmetic and uninterpreted functions (QF-UFLIA [Barrett et al. 2010]):

```
type Nat = { v:Integer | 0 ≤ v }
```

Specification & Verification Throughout this section, we will use the textbook Fibonacci function to demonstrate the proof features we add to LIQUID HASKELL, which we type as follows.

```
fib :: Nat → Nat
fib 0 = 0
fib 1 = 1
fib n = fib (n-1) + fib (n-2)
```

To ensure termination, the input type's refinement specifies a *pre-condition* that the parameter must be `Nat`. The output type's refinement specifies a *post-condition* that the result is also a `Nat`. Refinement type checking can automatically verify that if `fib` is invoked with a non-negative `Integer`, then it terminates and yields a non-negative `Integer`.

Propositions We can define a data type representing propositions as an alias for unit:

```
type Prop = ()
```

which can be *refined* with propositions about the code, *e.g.* that 2 + 2 equals 4

```
type Plus_2_2 = { v: Prop | 2 + 2 = 4 }
```

For simplicity, we abbreviate the above to **type** `Plus_2_2 = { 2 + 2 = 4 }`.

Universal & Existential Propositions Refinements encode universally-quantified propositions as *dependent function* types of the form:

```
type Plus_com = x:Integer → y:Integer → { x + y = y + x }
```

As x and y refer to arbitrary inputs, any inhabitant of the above type is a proof that `Integer` addition commutes. Refinements encode existential quantification via *dependent pairs* of the form:

```
type Nat_up = n:Nat → (m::Integer, {n < m})
```

The notation (m :: t, t') describes dependent pairs where the name m for the first element can appear inside refinements for the second element. Thus, `Nat_up` states the proposition that for every natural number n *there exists* value that is larger than n.

While quantifiers cannot appear directly inside the refinements, dependent functions and pairs allow us to specify quantified propositions. One limitation of this encoding is that quantifiers cannot exist inside logical connectives (like ∧ and ∨). In this paper, we present how to encode logical connectives using data types, *e.g.* conjunction as product and disjunction as a union, and show how to specify arbitrary higher-order logic (HOL) propositions using refinement types and how to verify those propositions using refinement type checking (§ 3).

Proofs We *prove* the above propositions by writing Haskell programs, for example

```
plus_2_2 :: Plus_2_2    plus_com :: Plus_com    nat_up :: Nat_up
plus_2_2 = ()           plus_com = \x y → ()    nat_up = \n → (n+1,())
```

Standard refinement typing reduces the above to the respective *verification conditions* (VCs)

$$true \Rightarrow 2 + 2 = 4 \qquad \forall\, x,\, y\, .\, true \Rightarrow x + y = y + x \qquad \forall\, n\, .n < n + 1$$

which are easily deemed valid by the SMT solver, allowing us to prove the respective propositions.

A Note on Bottom: Readers familiar with Haskell's semantics may be concerned that "bottom", which inhabits all types, makes our proofs suspect. Fortunately, as described in Vazou et al. [2014],

Liquid Haskell ensures that all terms with non-trivial refinements provably terminate and evaluate to (non-bottom) values, which makes our proofs sound.

2.2 Refinement Reflection

Suppose we wish to prove properties about the fib function, *e.g.* that {fib 2 = 1}. Standard refinement type checking runs into two problems. First, for decidability and soundness, *arbitrary* user-defined functions cannot belong in the refinement logic, *i.e.* we cannot *refer* to fib in a refinement. Second, the only specification that a refinement type checker has about fib is its type Nat → Nat which is too weak to verify {fib 2 = 1}. To address both problems, we **reflect** fib into the logic which sets the three steps of refinement reflection in motion.

Step 1: Definition The annotation creates an *uninterpreted function* fib :: Integer → Integer in the refinement logic. By uninterpreted, we mean that the logical fib is *not* connected to the program function fib; in the logic, fib only satisfies the *congruence axiom* $\forall n, m. \ n = m \ \Rightarrow$ fib $n = $ fib m. On its own, the uninterpreted function is not terribly useful: we cannot check {fib 2 = 1} as the SMT solver *cannot* prove the following VC (which requires reasoning about the *definition* of fib)

$$true \Rightarrow \text{fib } 2 = 1$$

Step 2: Reflection In the next key step, we reflect the *definition* of fib into its refinement type by automatically strengthening the user defined type for fib to:

```
fib :: n:Nat → { v:Nat | v = fib n && fibP n }
```

where fibP is an alias for a refinement *automatically derived* from the function's definition:

```
fibP n = n == 0 ⇒ fib n = 0
       ∧ n == 1 ⇒ fib n = 1
       ∧ n >= 1 ⇒ fib n = fib (n-1) + fib (n-2)
```

Step 3: Application With the reflected refinement type, each application of fib in the code automatically *unfolds* the definition of fib *once* in the logic. We prove {fib 2 = 1} by:

```
pf_fib2 :: { fib 2 = 1 }
pf_fib2 = let { t0 = fib 0; t1 = fib 1; t2 = fib 2 } in ()
```

We write in bold red, **f**, to highlight places where the unfolding of f's definition is important. Via refinement typing, the above yields the following VC that is discharged by SMT, even though fib is uninterpreted:

$$((\text{fibP } 0) \land (\text{fibP } 1) \land (\text{fibP } 2)) \Rightarrow (\text{fib } 2 = 1)$$

Note that the verification of pf_fib2 relies merely on the fact that fib is applied to (*i.e.* unfolded at) 0, 1 and 2. The SMT solver automatically *combines* the facts, once they are in the antecedent. The following is also verified:

```
pf_fib2' :: { fib 2 = 1 }
pf_fib2' = [ fib 0, fib 1, fib 2 ]
```

In the next subsection, we will continue to use explicit, step-by-step proofs as above, but we introduce additional tools for proof composition. Then, in § 2.4 we will eliminate unnecessary details in such proofs, using *Proof by Logical Evaluation* (PLE) for automation.

2.3 Equational Proofs

We can structure proofs to follow the style of *calculational* or *equational* reasoning popularized in classic texts [Bird 1989; Dijkstra 1976] and implemented AGDA [Mu et al. 2009] and DAFNY [Leino and Polikarpova 2016]. To this end, we have developed a library of proof combinators that permits reasoning about equalities and linear arithmetic.

"Equation" Combinators We equip LIQUID HASKELL with a family of equation combinators, $\odot$, for logical operators in the theory QF-UFLIA, $\odot \in \{=, \neq, \leq, <, \geq, >\}$. (In Haskell code, to avoid collisions with existing operators, we further append a period "." to these operators, so that "=" becomes "====" instead.) The refinement type of $\odot$ *requires* that $x \odot y$ holds and then *ensures* that the returned value is equal to x. For example, we define ==== as:

```
(====) :: x:a → y:{ a | x = y} → { v:a | v = x }
x ==== _ = x
```

and use it to write the following "equational" proof:

```
fib2_1 :: { fib 2 = 1 }
fib2_1 = fib 2 ==== fib 1 + fib 0 ==== 1 ** QED
```

where ** QED constructs "proof terms" by "casting" expressions to Prop in a post-fix fashion.

"Because" Combinators Often, we need to compose lemmas into larger theorems. For example, to prove fib 3 = 2 we may wish to reuse fib2_1 as a lemma. We do so with a "because" combinator:

```
(∵) :: (Prop → a) → Prop → a
f ∵ y = f y
```

The operator is simply an alias for function application that lets us write x $\odot$ y ∵ p. We use the because combinator to prove that fib 3 = 2.

```
fib3_2 :: { fib 3 = 2 }
fib3_2 = fib 3 ==== fib 2 + fib 1 ==== 2 ∵ fib2_1 ** QED
```

Here fib 2 is not important to unfold, because fib2_1 already provides the same information.

Arithmetic and Ordering Next, lets see how we can use arithmetic and ordering to prove that fib is (locally) increasing, *i.e.* for all n, fib $n \leq$ fib $(n + 1)$.

```
type Up f = n:Nat → {f n ≤ f (n + 1)}

fibUp   :: Up fib
fibUp 0 = fib 0 <. fib 1                          ** QED
fibUp 1 = fib 1 ≤. fib 1 + fib 0      ==== fib 2     ** QED
fibUp n = fib n ≤. fib n + fib (n-1) ==== fib (n+1) ** QED
```

Case Splitting The proof fibUp works by splitting cases on the value of n. In the *base* cases 0 and 1, we simply assert the relevant inequalities. These are verified as the reflected refinement unfolds the definition of fib at those inputs. The derived VCs are (automatically) proved as the SMT solver concludes $0 < 1$ and $1 + 0 \leq 1$ respectively. When n is greater than two, fib n is unfolded to fib (n−1) + fib (n−2), which, as fib (n−2) is non-negative, completes the proof.

Induction & Higher Order Reasoning Refinement reflection smoothly accomodates induction and higher-order reasoning. For example, lets prove that every function f that increases locally (*i.e.* f z $\leq$ f (z+1) for all z) also increases globally (*i.e.* f x $\leq$ f y for all x < y)

```
type Mono = f:(Nat → Integer) → Up f → x:_ → y:{x < y} → {f x ≤ f y}

fMono :: Mono / [y]
fMono f up x y
  | x+1 == y = f x ≤. f (x+1) ∵ up x ≤. f y                              ** QED
  | x+1  < y = f x ≤. f (y-1) ∵ fMono f up x (y-1) ≤. f y ∵ up (y-1) ** QED
```

We prove the theorem by induction on y as specified by the annotation / [y] which states that y
is a well-founded termination metric that decreases at each recursive call [Vazou et al. 2014]. If
x+1 == y, then we call the fUp x proof argument. Otherwise, x+1 < y, and we use the induction
hypothesis *i.e.* apply fMono at y-1, after which transitivity of the less-than ordering finishes the
proof. We can *apply* the general fMono theorem to prove that fib increases monotonically:

```
fibMono :: n:Nat → m:{n < m} → {fib n ≤ fib m}
fibMono = fMono fib fibUp
```

2.4 Complete Verification: Automating Equational Reasoning

While equational proofs can be very elegant, writing them out can quickly get exhausting. Lets face
it: fib3_2 is doing rather a lot of work just to prove that fib 3 equals 2! Happily, the *calculational*
nature of such proofs allows us to develop the following proof search algorithm PLE that is inspired
by model checking [Clarke et al. 1992]:

- **Step 1: Guard Normal Form** First, as shown in the definition of fibP above, each reflected
  function is transformed into a *guard normal form* $\wedge_i p_i \Rightarrow f(\overline{x}) = b_i$ *i.e.* as a collection of
  guards p_i and their corresponding definition b_i.
- **Step 2: Unfolding** Second, given a VC of the form $\Phi \Rightarrow p$, we iteratively *unfold* function
  application terms in Φ and p by *instantiating* them with the definition corresponding to
  an *enabled* guard. For example, given a VC *true* $\Rightarrow$ fib 3 = 2, the guard $3 \geq 1$ is trivially
  enabled, *i.e.* is true, and hence we strengthen the hypothesis Φ with the equality fib 3 =
  fib $3 - 1 +$ fib $3 - 2$ corresponding to unfolding the definition of fib at 3.
- **Step 3: Fixpoint** Third, we repeat the above process until either the goal is proved or we
  have reached a fixpoint, *i.e.* no further unfolding is enabled. For example, the above fixpoint
  computation unfolds the definition of fib at 2 and 1 and 0 and then stops as no further
  guards are enabled.

Automatic Equational Reasoning In § 6 we formalize a notion of *equational proof* and show
that the proof search procedure PLE enjoys two key properties. First, that it is guaranteed to find
an equational proof if one exists. Second, that under certain conditions readily met in practice, it is
guaranteed to terminate. These two properties allow us to use PLE to predictably automate proofs:
the programmer needs only supply the relevant induction hypotheses or helper lemma applications.
The remaining long chains of calculations are performed automatically via SMT-based PLE. To wit,
with complete proof search, the above proofs shrink to:

```
fib3_2 :: {fib 3 = 2}    fibUp  :: Up fib    fMono :: Mono / [y]
fib3_2 = ()              fibUp 0 = ()        fMono f up x y
                         fibUp 1 = ()          | x+1 == y = up x
                         fibUp n = ()          | x+1  < y = up (y-1) &&& fMono up x (y-1)
```

where the proof combinators x &&& y = () simply inserts the two proof arguments in the VC
environment.

```
app_assoc :: AppendAssoc                    app_assoc            :: AppendAssoc
app_assoc [] ys zs                          app_assoc []     ys zs = ()
  = ([] ++ ys) ++ zs                         app_assoc (x:xs) ys zs = app_assoc xs ys zs
  ==== ys ++ zs
  ==== [] ++ (ys ++ zs)        ** QED       app_right_id          :: AppendNilId
app_assoc (x:xs) ys zs                       app_right_id []       = ()
  = ((x : xs) ++  ys)  ++ zs                 app_right_id (x:xs)   = app_right_id xs
  ==== (x : (xs ++  ys)) ++ zs
  ====  x :((xs ++  ys)  ++ zs)             map_fusion            :: MapFusion
     ∵ app_assoc xs ys zs                    map_fusion f g []     = ()
  ====  x : (xs ++ (ys   ++ zs))            map_fusion f g (x:xs) = map_fusion f g xs
  ==== (x : xs) ++ (ys   ++ zs) ** QED
```

Fig. 1. (L) Equational proof of append associativity, (R) PLE proof, also of append-id and map-fusion.

PLE vs. Axiomatization Existing SMT based verifiers like DAFNY [Leino 2010] and F* [Swamy et al. 2016] use the classical *axiomatic* approach to verifying assertions over user-defined functions fib. In these systems, the function is encoded in the logic as a universally quantified formula (or axiom): $\forall n.$ fibP n after which the SMT solver may instantiate the above axiom at 2, 1 and 0 in order to automatically prove {fib 3 = 2}.

The automation offered by axioms is a bit of a devil's bargain, as axioms render VC checking *undecidable*, and in practice automatic axiom instantation can easily lead to infinite "matching loops". For example, the existence of a term fib n in a VC can trigger the above axiom, which may then produce the terms fib $(n - 1)$ and fib $(n - 2)$, which may then recursively give rise to further instantiations *ad infinitum*. To prevent matching loops an expert must carefully craft "triggers" or, alternatively provide a "fuel" parameter [Amin et al. 2014] that bounds the depth of instantiation. Both these approaches ensure termination, but can cause the axiom to not be instantiated at the right places, thereby rendering the VC checking *incomplete*. The incompleteness is illustrated by the following example from the DAFNY benchmark suite [Leino 2016]

```
pos n | n < 0    = 0            test  :: y:{y > 5} → {pos n = 3 + pos (n-3)}
      | otherwise = 1 + pos (n-1)    test _ = ()
```

DAFNY (and F*'s) fuel-based approach fails to check the above, when the fuel value is less than 3. One could simply raise-the-fuel-and-try-again but at what point does the user know when to stop? In contrast, PLE (1) does not require any fuel parameter, (2) is able to automatically perform the required unfolding to verify this example, *and* (3) is guaranteed to terminate.

2.5 Case Study: Laws for Lists

Reflection and PLE are not limited to integers. We end the overview by showing how they verify textbook properties of lists equipped with append (++) and map functions:

```
reflect (++) :: [a] → [a] → [a]    reflect map :: (a → b) → [a] → [b]
[]     ++ ys = ys                  map f []    = []
(x:xs) ++ ys = x : (xs ++ ys)      map f (x:xs) = f x : map f xs
```

In § 5.1 we will describe how the reflection mechanism illustrated via fibP is extended to account for ADTs using SMT-decidable selection and projection operations, which reflect the definition of ++ into the refinement: if isNil xs then ys else sel1 xs : (sel2 xs ++ ys) Note that LIQUID

HASKELL automatically checks that ++ and map are total [Vazou et al. 2014], which lets us safely **reflect** them into the refinement logic.

Laws We can specify various laws about lists with refinement types. For example, the below laws state that (1) appending to the right is an *identity* operation, (2) appending is an *associative* operation, and (3) map *distributes* over function composition:

```
type AppendNilId = xs:_ → { xs ++ [] = xs }
type AppendAssoc = xs:_ → ys:_ → zs:_ → { xs ++ (ys ++ zs) = (xs ++ ys) ++ zs }
type MapFusion   = f:_  → g:_  → xs:_ → { map (f . g) xs = map (f . map g) xs }
```

Proofs On the right in Figure 1 we show the proofs of these laws using PLE, which should be compared to the classical equational proof *e.g.* [Wadler 1987] shown on the left. With PLE, the user need only provide the high-level structure — the case splits and invocations of the induction hypotheses — after which PLE automatically completes the rest of the equational proof. Thus using SMT-based PLE, app_assoc shrinks down to its essence: an induction over the list xs. The difference is even more stark with map_fusion whose full equational proof we don't show, as it is twice as long.

PLE *vs. Normalization* The proofs in Figure 1 may remind readers familiar with Type-Theory based proof assistants (*e.g.* COQ or AGDA) of the notions of *type-level normalization* and *rewriting* that permit similar proofs in those systems. While our approach of PLE is inspired by the idea of type level computation, it differs from it in two significant ways. First, from a *theoretical* point of view, SMT logics are not equipped with any notion of computation, normalization, canonicity or rewriting. Instead, our PLE algorithm shows how to *emulate* (the spirit!) of those ideas by asserting equalities corresponding to function definitions. Second, from a *practical* perspective, the combination of PLE and (decidable) SMT-based theory reasoning can greatly simplify proofs. For example, consider the swap function from a COQ textbook [Appel 2016]:

```
swap :: [Integer] → [Integer]
swap (x1:x2:xs) | x1 > x2   = x2:x1:x2
                | otherwise = x1:x2:xs
swap xs                     = xs
```

In Figure 2 we show three proofs that swap is idempotent: Appel's proof using COQ (simplified by the use of a hint database and the arithmetic tacic omega), and variant in AGDA and finally the PLE proof. It is readily apparent that PLE's combination of proof search working hand-in-glove with SMT-based theory reasoning makes proving the result relatively trivial. Of course, the decades-worth of tactics, libraries and proof scripts available in AGDA, COQ, ISABELLE etc. enable large scale proof engineering that is well beyond what is currently feasible with our approach. We merely use this example it to illustrate that reflection and SMT-based proof search bring powerful, complete new tools for specification and verification.

3 COMPLETE SPECIFICATION

In this section, we show that the restriction in λ^R for the refinement language to be quantifier free—crucial for SMT-decidable type checking—does not pose any expressiveness restrictions to our specifications. Instead, quantified specifications can be naturally encoded using λ-abstractions and dependent pairs to encode universal and existential quantification respectively [Howard 1980]. In this section we assume the requirement that all the expressions are total.

COQ	AGDA	PBE

```
Theorem swap_idemp:
  ∀ al, swap (swap al) = swap al.
Proof. intros.
      destruct al as [ | a al].
      simpl. reflexivity.
      destruct al as [ | b al].
      simpl. reflexivity.
      simpl. bdestruct (b <? a).
  *
      simpl. bdestruct (a <? b).
      omega. reflexivity.
  *
      simpl. bdestruct (b <? a).
      omega. reflexivity.   Qed.
```

```
swap_idemp : ∀ xs → swap (swap xs) ≡ swap xs
swap_idemp [] = P.refl
swap_idemp (x₁ :: []) = P.refl
swap_idemp (x₁ :: x₂ :: xs) with x₁ ≤? x₂
swap_idemp (x₁ :: x₂ :: xs) | yes p with x₂ ≤? x₁
... | yes q rewrite antisym p q = P.refl
... | no ¬q = P.refl
swap_idemp (x₁ :: x₂ :: xs) | no ¬p with x₁ ≤? x₂
... | yes q = ⊥−elim (¬p q)
... | no ¬q = P.refl
```

```
swap_idemp ::
  xs:_ -> {swap (swap xs) = swap xs}

swap_idemp (x1:x2:xs)
  | x1 > x2   = ()
  | otherwise = ()
swap_idemp xs = ()
```

Fig. 2. Proofs that swap is idempotent with COQ, AGDA and PLE.

$$
\begin{aligned}
\phi \quad ::= & \qquad\qquad\qquad \textit{Formulas:} \\
& \{e\} & \textit{native terms as in Figure 6, with } \mathtt{True}, \mathtt{False} \in e \\
| \quad & \phi_1 \to \phi_2 & \textit{implication: } \phi_1 \Rightarrow \phi_2 \\
| \quad & \phi \to \{\mathtt{False}\} & \textit{negation: } !\phi \\
| \quad & \mathsf{PAnd}\ \phi_1\ \phi_2 & \textit{conjunction: } \phi_1 \wedge \phi_2 \\
| \quad & \mathsf{POr}\ \phi_1\ \phi_2 & \textit{disjunction: } \phi_1 \vee \phi_2 \\
| \quad & x : \tau \to \phi & \textit{forall: } \forall x.\phi \\
| \quad & (x :: \tau, \phi) & \textit{exists: } \exists x.\phi
\end{aligned}
$$

Fig. 3. Encoding of higher-order specifications in using quantifier-free refinement types. $\{e\}$ simplifies $\{v : \mathsf{Prop} \mid e\}$. Function binders are not represented in negation and implication where they are not relevant.

3.1 The specification language

Figure 3 summarizes the encoding of quantified specifications using a propositional (*i.e.* quantifier-free) refinements of λ^R.

Native terms Native terms consist of all (quantifier-free) expressions of λ^R.

Boolean connectives Implication $\phi_1 \Rightarrow \phi_2$ is encoded as a function from the proof of ϕ_1 to the proof of ϕ_2. Negation is encoded as an implication where the consequent is False. Conjunction $\phi_1 \wedge \phi_2$ is encoded with the data type PAnd that contains the proofs of the two conjuncts. and disjunction $\phi_1 \vee \phi_2$ is encoded with the sum type POr that contains the proofs of one of the two disjuncts.

```
data PAnd a b = PAnd a b
data POr  a b = POrL a | POrR b
```

Quantifiers Universal quantification $\forall x.\phi$ is encoded and introduced as lambda abstraction $x : a \to \phi$ and eliminated by function application. Existential quantification $\exists x.\phi$ is encoded as a dependent pair (x::a, ϕ) that contains x and a proof of a formula that depends on x. Even though λ^R does not have explicit syntax for dependent pairs, dependent pairs can be implemented using abstract refinement types [Vazou et al. 2013]. without adding extra complexity to the system.

$$\frac{\Gamma \vdash e_1 : \phi_1 \quad \Gamma \vdash e_2 : \phi_1}{\Gamma \vdash \mathsf{PAnd}\ e_1\ e_2 : \mathsf{PAnd}\ \phi_1\ \phi_2} \wedge\text{-I}$$

$$\frac{\Gamma \vdash \mathsf{PAnd}\ e_1\ e_2 : \mathsf{PAnd}\ \phi_1\ \phi_2}{\Gamma \vdash e_1 : \phi_1} \wedge\text{-E-L} \quad \frac{\Gamma \vdash \mathsf{PAnd}\ e_1\ e_2 : \mathsf{PAnd}\ \phi_1\ \phi_2}{\Gamma \vdash e_2 : \phi_2} \wedge\text{-E-R}$$

$$\frac{\Gamma \vdash e_1 : \phi_1}{\Gamma \vdash \mathsf{POrL}\ e_1 : \mathsf{POr}\ \phi_1\ \phi_2} \vee\text{-I-L} \quad \frac{\Gamma \vdash e_2 : \phi_2}{\Gamma \vdash \mathsf{POrR}\ e_2 : \mathsf{POr}\ \phi_1\ \phi_2} \vee\text{-I-R}$$

$$\frac{\Gamma \vdash e : \mathsf{POr}\ \phi_1\ \phi_2 \quad \Gamma, x_1{:}\phi_1 \vdash e_1 : \phi \quad \Gamma, x_2{:}\phi_2 \vdash e_2 : \phi}{\Gamma \vdash \mathsf{case}\ x = e\ \mathsf{of}\ \{\mathsf{POrL}\ x_1 \to e_1; \mathsf{POrR}\ x_2 \to e_2\} : \phi} \vee\text{-E}$$

$$\frac{\Gamma, x{:}\phi_x \vdash e : \phi}{\Gamma \vdash \lambda x.e : \phi_x \to \phi} \Rightarrow\text{-I} \quad \frac{\Gamma \vdash e : \phi_x \to \phi \quad \Gamma \vdash e_x : \phi_x}{\Gamma \vdash e\ e_x : \phi} \Rightarrow\text{-E}$$

$$\frac{\Gamma \vdash e_x : \phi \quad \Gamma \vdash e : \phi \to \{\mathsf{False}\}}{\Gamma \vdash e\ e_x : \{\mathsf{False}\}} \text{!-E} \quad \frac{\Gamma, x{:}\phi \vdash e : \{\mathsf{False}\}}{\Gamma \vdash \lambda x.e : \phi \to \{\mathsf{False}\}} \text{!-I}$$

$$\frac{\Gamma \vdash e_x : \tau \quad \Gamma \vdash e : (x : \tau \to \phi)}{\Gamma \vdash e\ e_x : \phi[x/e_x]} \forall\text{E} \quad \frac{\Gamma, x{:}\tau \vdash e : \phi}{\Gamma \vdash \lambda x.e : (x : \tau \to \phi)} \forall\text{I}$$

$$\frac{\Gamma \vdash e : (x :: \tau, \phi_x) \quad \Gamma, x{:}\tau, y{:}\phi_x \vdash e' : \phi}{\Gamma \vdash \mathsf{case}\ x = e\ \mathsf{of}\ \{(x, y) \to e'\} : \phi} \exists\text{E} \quad \frac{\Gamma \vdash e_x : \tau \quad \Gamma, x{:}\tau \vdash e : \phi}{\Gamma \vdash (e_x, e) : (x :: \tau, \phi[x/e_x])} \exists\text{I}$$

Fig. 4. Natural deduction [Gentzen 1935] rules for λ^R.

3.2 Proof Terms: Proofs via Natural Deduction

We defined ϕ to be both a proposition and a refinement type. We connect these two meanings of ϕ by proving that if there exist an expression with refinement type ϕ, the the proposition ϕ is valid.

THEOREM 3.1 (VALIDITY). *If $\emptyset; \emptyset \vdash e : \phi$ then ϕ is valid.*

We prove the theorem by soundness of λ^R (Theorem 4.1) and the fact that if the set $(\!|\phi|\!)$ is not empty then ϕ is valid, which is an implication of the Curry-Howard isomorphism [Howard 1980].

But how does one construct a proof term e for a formula ϕ? For this construction we can use Gentzen's natural deduction system. Figure 4 maps each natural deduction derivation rule from [Gentzen 1935] to a typing judgment accepted by λ^R. To get the rules for natural deduction one should read $\Gamma \vdash e : \phi$ as "ϕ is provable under the assumptions of Γ" and $\Gamma \vdash e : \tau$ as "e is a term". Let $\Gamma \vdash_{ND} \phi$ denote Gentzen's natural deduction rules, then for example, conjunction and universal elimination derive from the respective rules of Figure 4 as follows.

$$\frac{\Gamma \vdash_{ND} \phi_1 \vee \phi_2 \quad \Gamma, \phi_1 \vdash_{ND} \phi \quad \Gamma, \phi_2 \vdash_{ND} \phi}{\Gamma \vdash_{ND} \phi} \vee\text{-E} \quad \frac{\Gamma \vdash_{ND} e_x\ \mathsf{term} \quad \Gamma \vdash_{ND} \forall x.\phi}{\Gamma \vdash_{ND} \phi[x/e_x]} \forall\text{E}$$

Since Figure 4 maps directly natural deduction rules to derivations that are accepted by λ^R, we conclude that if there is a natural deduction derivation for a proposition ϕ, then there exists a λ^R term that proves this formula.

THEOREM 3.2. *If $\emptyset \vdash_{ND} \phi$ then we can construct an e so that $\emptyset; \emptyset \vdash e : \phi$*

3.3 Examples

Next illustrate examples of proofs for quantified propositions. These propositions range from textbook logical properties to properties over user-defined domains (here lists), and even induction on integers.

$$\cfrac{\cfrac{p:\phi_p, y:\tau_y \vdash p : \phi_p \qquad \cfrac{\cfrac{p:\phi_p, y:\tau_y, x:t_x, p_x:\phi_x \vdash p_x : \phi_x}{p:\phi_p, y:\tau_y, x:t_x, p_x:\phi_x \vdash p_x\ y : f\ x\ y}\quad\text{∀-E}}{\cfrac{p:\phi_p, y:\tau_y, x:t_x, p_x:\phi_x \vdash y : \tau_y}{}}}{\cfrac{p:\phi_p, y:\tau_y \vdash \text{case } p \text{ of } \{(x, p_x) \to (x, p_x\ y)\} : \exists x.(f\ x\ y)}{\cfrac{p:\phi_p \vdash \lambda y.\text{case } p \text{ of } \{(x, p_x) \to (x, p_x\ y)\} : \forall y.\exists x.(f\ x\ y)}{\emptyset \vdash \lambda p\ y.\text{case } p \text{ of } \{(x, p_x) \to (x, p_x\ y)\} : (\exists x.\forall y.(f\ x\ y)) \Rightarrow (\forall y.\exists x.(f\ x\ y))}\ \Rightarrow\text{-I}}\ \text{∀-I}}\ \exists\text{-E}}$$

Fig. 5. Proof of $(\exists x.\forall y.(f\ x\ y)) \Rightarrow (\forall y.\exists x.(f\ x\ y))$ where $\phi_p \equiv \exists x.\forall y.(f\ x\ y)$, $\phi_x \equiv \forall y.(f\ x\ y)$.

3.3.1 Natural Deduction as Type Derivation. To see the mapping from natural deduction to typing rules in action Figure 5 using typing judgments to express the the Gentzen's proof of the proposition

$$\phi \equiv (\exists x.\forall y.(f\ x\ y)) \Rightarrow (\forall y.\exists x.(f\ x\ y))$$

Reading bottom up the derivation provides a proof of ϕ, while reading top down it constructs the proof term of the formula to be $\lambda p\ y.\text{case } p \text{ of } \{(x, p_x) \to (x, p_x\ y)\}$.

3.3.2 Logical Properties. Next, we prove distribution of existentials over disjunction:

$$\phi \equiv \exists x.(p\ x \lor q\ x)) \Rightarrow ((\exists x.p\ x) \lor (\exists x.q\ x))$$

The proof is a λ^R program (here expressed using Haskell). The specification of this property requires nesting existentials inside disjunctions and vice versa. The proof proceeds by existential case splitting and introduction:

```
existsOrDistr :: p:(a → Bool) → g:(a → Bool)
               → (x::a, POr {p x} {q x})
               → POr (x::a, {p x}) (x::a, {q x})
existsOrDistr _ _ (x,POrL px) = POrL  (x,px)
existsOrDistr _ _ (x,POrR qx) = POrR (x,qx)
```

Similarly, we distribute foralls over conjunction:

$$\phi \equiv \forall x.(p\ x \land q\ x)) \Rightarrow ((\forall x.p\ x) \land (\forall x.q\ x))$$

The specification of the conclusion now requires nesting universal quantification over conjunctions. This requirement leads to a proof term that performs λ-abstraction and case spitting inside the PAnd data constructor.

```
forallAndDistr :: p:(a → Bool) → q:(a → Bool)
               → (x:a → PAnd {p x} {q x})
               → PAnd (x:a → {p x}) (x:a → {q x})
forallAndDistr _ _ andx
  = PAnd (\x → case andx x of PAnd px _ → px)
         (\x → case andx x of PAnd _ qx → qx)
```

3.3.3 Properties on user specified domains. Since native terms are drawn from user-defined decidable logics, the terms in ϕ can talk about properties of data types, like lists. As an we prove that forall lists xs if there exists a list ys so that xs == ys ++ ys then xs has even length.

$$\phi \equiv (\forall xs.\exists ys.xs = ys\text{++}ys) \Rightarrow (\exists n.\text{length } xs = n + n)$$

The proof proceeds by existential elimination and introduction, and by invocation of the `lenAppend` lemma.

```
even_lists :: xs:[a] → (ys::[a], {xs == ys ++ ys})
             → (n::Int, {length xs == n + n})
even_lists xs (ys,pf) = (length ys, lenAppend ys ys &&& pf)


lenAppend :: xs:[a] → ys:[a]
             → {length (xs ++ ys) == length xs + length ys}
lenAppend []     _   = ()
lenAppend (x:xs) ys = lenAppend xs ys
```

The `lenAppend` lemma is proven by induction on the input list, while PBE is used to simplify the trivial unfoldings.

3.3.4   Induction on Natural Numbers. Finally, we use Liquid Haskell to specify and verify induction on natural numbers.

$$\phi \equiv (p\ 0 \wedge (\forall n.p\ (n-1) \Rightarrow p\ n)) \Rightarrow \forall n.p\ n$$

```
natInd :: p:(Int → Bool) → PAnd {p 0} (n:Int → {p (n-1)} → {p n})
        → n:Nat → {p n}
natInd p (PAnd p0 pn) n
  | n == 0    = p0
  | otherwise = pn n (natInd p (PAnd p0 pn) (n-1))
```

The proof proceeds by induction (*e.g.* case splitting). In the base case, `n == 0`, the proof calls the left conjunct. Otherwise, `0 < n`, the proof calls the right conjuct instantiated on the correct argument `n` and assuming the inductive hypothesis.

3.4   Consequences

Our discussion for far is not novel, it is merely an application of Curry-Howard correspondence to our system. What is novel is the evidence that Curry-Howard correspondence actually apply to λ^R which leads to two major contributions.

Fist, we show that natural deduction reasoning can smoothly co-exists with SMT verification to automate both the quantifier-free portions of the proof. For each (quantifier-free) propositional connective of ϕ, we can define a refinement and a reification operator that maps propositions to a refinement type and back. For instance, we refine conjunction by performing case analysis on the `PAnd {b1} {b2}` constructor therefore bringing both the conjuncts into the environment leading to the trivial VC $b1 \wedge b2 \Rightarrow b1 \wedge b2$ that can be trivially discharged by the SMT-solver.

```
andRefine :: b1:Bool → b2:Bool → PAnd {b1} {b2} → {b1 && b2}
andRefine _ _ (PAnd b1 b2) = ()
```

We reify conjunction by using $\phi_1 \wedge \phi_2$ as a proof for each conjunct ϕ_1 and ϕ_2.

```
andReify :: b1:Bool → b2:Bool → {b1 && b2} → PAnd {b1} {b2}
andReify _ _ b = PAnd b b
```

Second, since λ^R is implemented in the legacy programming language Liquid Haskell we show for first time how natural deduction proofs are encoded in Liquid Haskell which both sets clearer bounds for the expressiveness of the language and provides a pleasant implementation of natural deduction that can be used for pedagogical purposes.

$$
\begin{array}{rrcl}
\textbf{\textit{Operators}} & \odot & ::= & = \mid < \\
\textbf{\textit{Constants}} & c & ::= & \wedge \mid\, ! \mid \odot \mid +, -, \dots \\
& & \mid & \mathtt{True} \mid \mathtt{False} \mid 0, \pm 1, \dots \\
\textbf{\textit{Values}} & w & ::= & c \mid \lambda x.e \mid D\,\overline{w} \\
\textbf{\textit{Expressions}} & e & ::= & w \mid x \mid e\,e \\
& & \mid & \mathtt{case}\ x = e\ \mathtt{of}\ \{D\,\overline{x} \to e\} \\
\textbf{\textit{Binders}} & b & ::= & e \mid \mathtt{let\ rec}\ x : \tau = b\ \mathtt{in}\ b \\
\textbf{\textit{Program}} & p & ::= & b \mid \mathtt{reflect}\ x : \tau = e\ \mathtt{in}\ p \\
\textbf{\textit{Basic Types}} & B & ::= & \mathtt{Int} \mid \mathtt{Bool} \mid T \\
\textbf{\textit{Ref. Types}} & \tau & ::= & \{v : B^{[\Downarrow]} \mid e\} \mid x{:}\tau_x \to \tau
\end{array}
$$

Fig. 6. **Syntax of** λ^R: a calculus with an undecidable type system

4 REFINEMENT REFLECTION: λ^R

We formalize refinement reflection in three steps. First, we develop a core calculus λ^R with an *undecidable* type system based on denotational semantics. We show how the soundness of the type system allows us to *prove theorems* using λ^R. Next, in § 5 we define a language λ^S that soundly approximates λ^R while enabling decidable SMT-based type checking. Finally, in § 6 we strengthen proof obligations with complete unfoldings of the reflected functions, to automate equational reasoning.

4.1 Syntax

Figure 6 summarizes the syntax of λ^R, which is essentially the calculus λ^U [Vazou et al. 2014] with explicit recursion and a special reflect binding to denote terms that are reflected into the refinement logic. The elements of λ^R are layered into primitive constants, values, expressions, binders and programs.

Constants The primitive constants of λ^R include primitive relations $\oplus$, here, the set $\{=, <\}$. Moreover, they include the primitive booleans $\mathtt{True}$, $\mathtt{False}$, integers -1, 0, 1, *etc.*, and logical operators $\wedge$, $\vee$, $!$, *etc.*.

Data Constructors Data constructors are special constants. For example, the data type $[Int]$, which represents finite lists of integers, has two data constructors: $[]$ (nil) and $:$ (cons).

Values & Expressions The values of λ^R include constants, λ-abstractions $\lambda x.e$, and fully applied data constructors D that wrap values. The expressions of λ^R include values, variables x, applications $e\,e$, and case expressions.

Binders & Programs A *binder* b is a series of possibly recursive let definitions, followed by an expression. A *program* p is a series of reflect definitions, each of which names a function that is reflected into the refinement logic, followed by a binder. The stratification of programs via binders is required so that arbitrary recursive definitions are allowed in the program but cannot be inserted into the logic via refinements or reflection. (We *can* allow non-recursive let binders in expressions e, but omit them for simplicity.)

4.2 Operational Semantics

We define $\hookrightarrow$ to be the small step, call-by-name β-reduction semantics for λ^R. We evaluate reflected terms as recursive let bindings, with extra termination-check constraints imposed by the type

system:

$$\text{reflect } x : \tau = e \text{ in } p \hookrightarrow \text{let rec } x : \tau = e \text{ in } p$$

We define $\hookrightarrow^\star$ to be the reflexive, transitive closure of $\hookrightarrow$. Moreover, we define $\approx_\beta$ to be the reflexive, symmetric, and transitive closure of $\hookrightarrow$.

Constants Application of a constant requires the argument be reduced to a value; in a single step, the expression is reduced to the output of the primitive constant operation, *i.e.* $c\,v \hookrightarrow \delta(c, v)$. For example, consider $=$, the primitive equality operator on integers. We have $\delta(=, n) \doteq =_n$ where $\delta(=_n, m)$ equals True iff m is the same as n.

Equality We assume that the equality operator is defined *for all* values, and, for functions, is defined as extensional equality. That is, for all f and f', $(f = f') \hookrightarrow$ True iff $\forall v.\ f\,v \approx_\beta f'\,v$. We assume source *terms* only contain implementable equalities over non-function types; while function extensional equality only appears in *refinements*.

4.3 Types

λ^R types include basic types, which are *refined* with predicates, and dependent function types. *Basic types* B comprise integers, booleans, and a family of data-types T (representing lists, trees *etc.*). For example, the data type $[Int]$ represents lists of integers. We refine basic types with predicates (boolean-valued expressions e) to obtain *basic refinement types* $\{v : B \mid e\}$. We use $\Downarrow$ to mark provably terminating computations and use refinements to ensure that if $e : \{v : B^\Downarrow \mid e'\}$, then e terminates. As discussed by [Vazou et al. \[2014\]](#) termination labels can be checked using refinement types and are used to ensure that refinements cannot diverge as required for soundness of type checking under lazy evaluation. Finally, we have dependent *function types* $x : \tau_x \to \tau$ where the input x has the type τ_x and the output τ may refer to the input binder x. We write B to abbreviate $\{v : B \mid \text{True}\}$, and $\tau_x \to \tau$ to abbreviate $x : \tau_x \to \tau$ if x does not appear in τ.

Denotations Each type τ *denotes* a set of expressions $[\![\tau]\!]$, that is defined via the operational semantics [[Knowles and Flanagan 2010](#)]. Let $\text{shape}(\tau)$ be the type we get if we erase all refinements from τ and $e : \text{shape}(\tau)$ be the standard typing relation for the typed lambda calculus. Then, we define the denotation of types as:

$$[\![\{x : B \mid r\}]\!] \doteq \{e \mid e : B, \text{ if } e \hookrightarrow^\star w \text{ then } r[x/w] \hookrightarrow^\star \text{True}\}$$

$$[\![\{x : B^\Downarrow \mid r\}]\!] \doteq [\![\{x : B \mid r\}]\!] \cap \{e \mid e \hookrightarrow^\star w\}$$

$$[\![x : \tau_x \to \tau]\!] \doteq \{e \mid e : \text{shape}(\tau_x \to \tau), \forall e_x \in [\![\tau_x]\!].\ (e\,e_x) \in [\![\tau[x/e_x]]\!]\}$$

Constants For each constant c we define its type $\text{prim}(c)$ such that $c \in [\![\text{prim}(c)]\!]$. For example,

$$\begin{aligned}
\text{prim}(3) &\doteq \{v : \text{Int}^\Downarrow \mid v = 3\} \\
\text{prim}(+) &\doteq x : \text{Int}^\Downarrow \to y : \text{Int}^\Downarrow \to \{v : \text{Int}^\Downarrow \mid v = x + y\} \\
\text{prim}(\leq) &\doteq x : \text{Int}^\Downarrow \to y : \text{Int}^\Downarrow \to \{v : \text{Bool}^\Downarrow \mid v \Leftrightarrow x \leq y\}
\end{aligned}$$

4.4 Refinement Reflection

The key idea in our work is to *strengthen* the output type of functions with a refinement that *reflects* the definition of the function in the logic. We do this by treating each reflect-binder (reflect $f : \tau = e$ in p) as a let rec-binder (let rec $f : \text{Reflect}(\tau, e) = e$ in p) during type checking (rule T-REFL in Figure [7](#)).

Reflection We write $\text{Reflect}(\tau, e)$ for the *reflection* of the term e into the type τ, defined by strengthening τ as:

$$\begin{aligned}
\text{Reflect}(\{v : B^{\Downarrow} \mid r\}, e) &\doteq \{v : B^{\Downarrow} \mid r \wedge v = e\} \\
\text{Reflect}(x : \tau_x \to \tau, \lambda x.e) &\doteq x : \tau_x \to \text{Reflect}(\tau, e)
\end{aligned}$$

As an example, recall from § 2 that the **reflect** fib strengthens the type of fib with the refinement fibP. That is, let the user specified type of fib be t_{fib} and the its definition be definition $\lambda n.e_{\text{fib}}$.

$$\begin{aligned}
t_{\text{fib}} &\doteq \{v : \text{Int} \mid 0 \le v\} \to \{v : \text{Int} \mid 0 \le v\} \\
e_{\text{fib}} &\doteq \text{case } x = n \le 1 \text{ of } \{\text{True} \to n; \text{False} \to \text{fib}(n-1) + \text{fib}(n-2)\}
\end{aligned}$$

Then, the reflected type of fib will be:

$$\text{Reflect}(t_{\text{fib}}, e_{\text{fib}}) = n : \{v : \text{Int} \mid 0 \le v\} \to \{v : \text{Int} \mid 0 \le v \wedge v = e_{\text{fib}}\}$$

Termination Checking We defined $\text{Reflect}(\cdot, \cdot)$ to be a *partial* function that only reflects provably terminating expressions, *i.e.* expressions whose result type is marked with $\Downarrow$. If a non-provably terminating function is reflected in an λ^R expression then type checking will fail (with a reflection type error in the implementation). This restriction is crucial for soundness, as diverging expressions can lead to inconsistencies. For example, reflecting the diverging f x = 1 + f x into the logic leads to an inconsistent system that is able to prove 0 = 1.

Automatic Reflection Reflection of λ^R expressions into the refinements happens automatically by the type system, not manually by the user. The user simply annotates a function f as reflect f. Then, the rule T-REFL in Figure 7 is used to type check the reflected function by strengthening the f's result via $\text{Reflect}(\cdot, \cdot)$. Finally, the rule T-LET is used to check that the automatically strengthened type of f satisfies f's implementation.

4.5 Typing Rules

Next, we present the type-checking rules of λ^R, as found in Figure 7.

Environments and Closing Substitutions A *type environment* Γ is a sequence of type bindings $x_1 : \tau_1, \ldots, x_n : \tau_n$. An environment denotes a set of *closing substitutions* θ which are sequences of expression bindings: $x_1 \mapsto e_1, \ldots, x_n \mapsto e_n$ such that:

$$[\![\Gamma]\!] \doteq \{\theta \mid \forall x : \tau \in \Gamma. \theta(x) \in [\![\theta \cdot \tau]\!]\}$$

where $\theta \cdot \tau$ applies a substitution to a type (and likewise $\theta \cdot p$, to a program).

A reflection environment R is a sequence that binds the names of the reflected functions with their definitions $f_1 \mapsto e_1, \ldots, f_n \mapsto e_n$. A reflection environment respects a type environment when all reflected functions satisfy their types:

$$\Gamma \models R \doteq \forall (f \mapsto e) \in R. \exists \tau. (f : \tau) \in \Gamma \wedge (\Gamma; R \vdash e : \tau)$$

Typing A judgment $\Gamma; R \vdash p : \tau$ states that the program p has the type τ in the environment Γ under the reflection environment R. That is, when the free variables in p are bound to expressions described by Γ, the program p will evaluate to a value described by τ.

Rules All but two of the rules are the standard refinement typing rules [Knowles and Flanagan 2010; Vazou et al. 2014] that moreover preserve the reflection environment R. First, rule T-REFL is used to extend the reflection environment with the binding of the function name with its definition $(f \mapsto e)$ and moreover strengthen the type of each reflected binder with its definition, as described previously in § 4.4. Second, rule T-EXACT strengthens the expression with a singleton type equating

Typing $\boxed{\Gamma; R \vdash p : \tau}$

$$\frac{x : \tau \in \Gamma}{\Gamma; R \vdash x : \tau} \text{ T-Var} \qquad \frac{}{\Gamma; R \vdash c : \mathsf{prim}(c)} \text{ T-Con} \qquad \frac{\Gamma; R \vdash p : \tau' \quad \Gamma; R \vdash \tau' \preceq \tau}{\Gamma; R \vdash p : \tau} \text{ T-Sub}$$

$$\frac{\Gamma; R \vdash e : \{v : B \mid e_r\}}{\Gamma; R \vdash e : \{v : B \mid e_r \wedge v = e\}} \text{ T-Exact} \qquad \frac{\Gamma, x : \tau_x; R \vdash e : \tau}{\Gamma; R \vdash \lambda x.e : x : \tau_x \rightarrow \tau} \text{ T-Fun}$$

$$\frac{\Gamma; R \vdash e_1 : (x : \tau_x \rightarrow \tau) \quad \Gamma; R \vdash e_2 : \tau_x}{\Gamma; R \vdash e_1 \, e_2 : \tau} \text{ T-App} \qquad \frac{\begin{array}{c} \Gamma, x : \tau_x; R \vdash b_x : \tau_x \quad \Gamma, x : \tau_x \vdash \tau_x \\ \Gamma, x : \tau_x; R \vdash b : \tau \quad \Gamma \vdash \tau \end{array}}{\Gamma; R \vdash \mathsf{let\ rec}\ x : \tau_x = b_x\ \mathsf{in}\ b : \tau} \text{ T-Let}$$

$$\frac{\begin{array}{c} \Gamma; R \vdash e : \{v : T \mid e_r\} \qquad \Gamma \vdash \tau \\ \forall i.\mathsf{prim}(D_i) = \overline{y_j : \tau_j} \rightarrow \{v : T \mid e_{r_i}\} \quad \Gamma, \overline{y_j : \tau_j}, x : \{v : T \mid e_r \wedge e_{r_i}\}; R \vdash e_i : \tau \end{array}}{\Gamma; R \vdash \mathsf{case}\ x = e\ \mathsf{of}\ \{D_i\ \overline{y_i} \rightarrow e_i\} : \tau} \text{ T-Case}$$

$$\frac{\Gamma; R, f \mapsto e \vdash \mathsf{let\ rec}\ f : \mathsf{Reflect}(\tau_f, e) = e\ \mathsf{in}\ p : \tau}{\Gamma; R \vdash \mathsf{reflect}\ f : \tau_f = e\ \mathsf{in}\ p : \tau} \text{ T-Refl}$$

Well Formedness $\boxed{\Gamma \vdash \tau}$

$$\frac{\Gamma, v : B; \emptyset \vdash e : \mathsf{Bool}^{\Downarrow}}{\Gamma \vdash \{v : B \mid e\}} \text{ WF-Base} \qquad \frac{\Gamma \vdash \tau_x \quad \Gamma, x : \tau_x \vdash \tau}{\Gamma \vdash x : \tau_x \rightarrow \tau} \text{ WF-Fun}$$

Subtyping $\boxed{\Gamma; R \vdash \tau_1 \preceq \tau_2}$

$$\frac{\forall \theta \in [\![\Gamma]\!].[\![\theta \cdot \{v : B \mid e_1\}]\!] \subseteq [\![\theta \cdot \{v : B \mid e_2\}]\!]}{\Gamma; R \vdash \{v : B \mid e_1\} \preceq \{v : B \mid e_2\}} \preceq\text{-Base-}\lambda^R$$

$$\frac{\Gamma; R \vdash \tau_x' \preceq \tau_x \quad \Gamma, x : \tau_x'; R \vdash \tau \preceq \tau'}{\Gamma; R \vdash x : \tau_x \rightarrow \tau \preceq x : \tau_x' \rightarrow \tau'} \preceq\text{-Fun}$$

Fig. 7. Typing of λ^R

the value and the expression (*i.e.* reflecting the expression in the type). This is a generalization of the "selfification" rules from [Knowles and Flanagan 2010; Ou et al. 2004] and is required to equate the reflected functions with their definitions. For example, the application $\mathtt{fib}\ 1$ is typed as $\{v : \mathtt{Int}^{\Downarrow} \mid \mathtt{fibP}\ v\ 1 \wedge v = \mathtt{fib}\ 1\}$ where the first conjunct comes from the (reflection-strengthened) output refinement of $\mathtt{fib}$ § 2 and the second comes from rule T-Exact.

Well-formedness A judgment $\Gamma \vdash \tau$ states that the refinement type τ is well-formed in the environment Γ. Following Vazou et al. [2014], τ is well-formed if all the refinements in τ are $\mathtt{Bool}$-typed, provably terminating expressions in Γ.

Subtyping A judgment $\Gamma; R \vdash \tau_1 \preceq \tau_2$ states that the type τ_1 is a subtype of τ_2 in the environments Γ and R. Informally, τ_1 is a subtype of τ_2 if, when the free variables of τ_1 and τ_2 are bound to expressions described by Γ, the denotation of τ_1 is *contained in* the denotation of τ_2. Subtyping of basic types reduces to denotational containment checking, shown in rule $\preceq$-Base-λ^R. That is, τ_1 is a subtype of τ_2 under Γ if for any closing substitution θ in the denotation of Γ, $[\![\theta \cdot \tau_1]\!]$ is contained in $[\![\theta \cdot \tau_2]\!]$.

Soundness Following λ^U [Vazou et al. 2014], in Supplementary-Material [2017] we prove that evaluation preserves typing and typing implies denotational inclusion.

$$
\begin{array}{rrcl}
\textit{\textbf{Predicates}} & p & ::= & p \bowtie p \mid \oplus_1 p \\
& & \mid & n \mid b \mid x \mid D \mid x\,\overline{p} \\
& & \mid & \text{if } p \text{ then } p \text{ else } p \\
\textit{\textbf{Integers}} & n & ::= & 0, -1, 1, \ldots \\
\textit{\textbf{Booleans}} & b & ::= & \text{True} \mid \text{False} \\
\textit{\textbf{Binary Ops}} & \bowtie & ::= & = \mid < \mid \wedge \mid + \mid - \mid \ldots \\
\textit{\textbf{Unary Ops}} & \oplus_1 & ::= & ! \mid \ldots \\
\textit{\textbf{Sort Args}} & s_a & ::= & \text{Int} \mid \text{Bool} \mid \mathsf{U} \mid \text{Fun } s_a\, s_a \\
\textit{\textbf{Sorts}} & s & ::= & s_a \mid s_a \to s
\end{array}
$$

Fig. 8. **Syntax of** λ^S

THEOREM 4.1. *[Soundness of λ^R]*

- **Denotations** *If* $\Gamma; R \vdash p : \tau$ *then* $\forall \theta \in [\![\Gamma]\!].\theta \cdot p \in [\![\theta \cdot \tau]\!]$.
- **Preservation** *If* $\emptyset; \emptyset \vdash p : \tau$ *and* $p \hookrightarrow^\star w$, *then* $\emptyset; \emptyset \vdash w : \tau$.

Theorem 4.1 lets us interpret well typed programs as proofs of propositions. For example, in § 2 we verified that the term fibUp proves $n{:}\text{Nat} \to \{\texttt{fib } n \le \texttt{fib } (n+1)\}$. Via soundness of λ^R, we get that for each valid input n, the result refinement is valid.

$$
\forall n.0 \le n \hookrightarrow^\star \text{True} \Rightarrow \texttt{fib } n \le \texttt{fib } (n+1) \hookrightarrow^\star \text{True}
$$

5 ALGORITHMIC CHECKING: λ^S

Next, we describe λ^S, a conservative, first order approximation of λ^R where higher-order features are approximated with uninterpreted functions and the undecidable type subsumption rule $\preceq$-BASE-λ^R is replaced with a decidable one $\preceq$-BASE-PBE, yielding an SMT-based algorithmic type system that is both sound and decidable.

Syntax: Terms & Sorts Figure 8 summarizes the syntax of λ^S, the *sorted* (SMT-) decidable logic of quantifier-free equality, uninterpreted functions and linear arithmetic (QF-EUFLIA) [Barrett et al. 2010; Nelson 1981]. The *terms* of λ^S include integers n, booleans b, variables x, data constructors D (encoded as constants), fully applied unary $\oplus_1$ and binary $\bowtie$ operators, and application $x\,\overline{p}$ of an uninterpreted function x. The *sorts* of λ^S include built-in integer Int and Bool for representing integers and booleans. The interpreted functions of λ^S, *e.g.* the logical constants = and <, have the function sort $s \to s$. Other functional values in λ^R, *e.g.* reflected λ^R functions and λ-expressions, are represented as first-order values with the uninterpreted sort Fun $s\ s$. The universal sort U represents all other values.

5.1 Transforming λ^R into λ^S

The judgment $\Gamma \vdash e \rightsquigarrow p$ states that a λ^R term e is transformed, under an environment Γ, into a λ^S term p. If $\Gamma \vdash e \rightsquigarrow p$ and Γ is clear from the context we write $\lfloor e \rfloor$ and $\lceil p \rceil$ to denote the translation from λ^R to λ^S and back. Most of the transformation rules are identity and can be found in [Supplementary-Material 2017]. Here we discuss the non-identity ones.

Embedding Types We embed λ^R types into λ^S sorts as:

$$
\begin{aligned}
(\!|\text{Int}|\!) &\doteq \text{Int} \\
(\!|\text{Bool}|\!) &\doteq \text{Bool} & (\!|\{v : B^{[\![\!]\!]} \mid e\}|\!) &\doteq (\!|B|\!) \\
(\!|T|\!) &\doteq \text{U} & (\!|x : \tau_x \to \tau|\!) &\doteq \text{Fun } (\!|\tau_x|\!) \, (\!|\tau|\!)
\end{aligned}
$$

Embedding Constants Elements shared on both λ^R and λ^S translate to themselves. These elements include booleans, integers, variables, binary and unary operators. SMT solvers do not support currying, and so in λ^S, all function symbols must be fully applied. Thus, we assume that all applications to primitive constants and data constructors are fully applied, *e.g.* by converting source terms like (+ 1) to (\z → z + 1).

Embedding Functions As λ^S is first-order, we embed λ-abstraction using the uninterpreted function lam.

$$
\frac{\Gamma, x : \tau_x \vdash e \rightsquigarrow p \qquad \Gamma; \Psi \vdash (\lambda x.e) : (x : \tau_x \to \tau)}{\Gamma \vdash \lambda x.e \rightsquigarrow \text{lam}_{(\!|\tau|\!)}^{(\!|\tau_x|\!)} \, x \, p}
$$

The term $\lambda x.e$ of type $\tau_x \to \tau$ is transformed to $\text{lam}_s^{s_x} \, x \, p$ of sort Fun $s_x \, s$, where s_x and s are respectively $(\!|\tau_x|\!)$ and $(\!|\tau|\!)$, $\text{lam}_s^{s_x}$ is a special uninterpreted function of sort $s_x \to s \to \text{Fun } s_x \, s$, and x of sort s_x and r of sort s are the embedding of the binder and body, respectively. As lam is an SMT-function, it *does not* create a binding for x. Instead, x is renamed to a *fresh* name pre-declared in the SMT logic.

Embedding Applications We embed applications via defunctionalization [Reynolds 1972] using the uninterpreted app:

$$
\frac{\Gamma \vdash e' \rightsquigarrow p' \qquad \Gamma \vdash e \rightsquigarrow p \qquad \Gamma; \Psi \vdash e : \tau_x \to \tau}{\Gamma \vdash e \, e' \rightsquigarrow \text{app}_{(\!|\tau|\!)}^{(\!|\tau_x|\!)} \, p \, p'}
$$

The term $e \, e'$, where e and e' have types $\tau_x \to \tau$ and τ_x, is transformed to $\text{app}_s^{s_x} \, p \, p' : s$ where s and s_x are $(\!|\tau|\!)$ and $(\!|\tau_x|\!)$, the $\text{app}_s^{s_x}$ is a special uninterpreted function of sort Fun $s_x \, s \to s_x \to s$, and p and p' are the respective translations of e and e'.

Embedding Data Types We data constructors to a predefined λ^S constant s_D of sort $(\!|\text{prim}(D)|\!)$. $\Gamma \vdash D \rightsquigarrow s_D$ For each datatype, we create reflected measures that *check* the top-level constructor and *select* their individual fields. For example, for lists, we create measures

```
isNil []    = True       isCons (x:xs) = True       sel1 (x:xs) = x
isNil (x:xs) = False      isCons []    = False       sel2 (x:xs) = xs
```

The above selectors can be modeled precisely in the refinement logic via SMT support for ADTs [Nelson 1981]. To generalize, let D_i be a data constructor such that $\text{prim}(D_i) \doteq \tau_{i,1} \to \cdots \to \tau_{i,n} \to \tau$ Then the *check function* is_{D_i} has the sort Fun $(\!|\tau|\!)$ Bool and the *select function* $\text{sel}_{D_{i,j}}$ has the sort Fun $(\!|\tau|\!) \, (\!|\tau_{i,j}|\!)$.

Embedding Case Expressions We translate case-expressions of λ^R into nested if terms in λ^S, by using the check functions in the guards and the select functions for the binders of each case.

$$
\frac{\Gamma \vdash e \rightsquigarrow p \qquad \Gamma \vdash e_i[\overline{y_i}/\overline{\text{sel}_{D_i} \, x}][x/e] \rightsquigarrow p_i}{\Gamma \vdash \text{case } x = e \text{ of } \{D_i \, \overline{y_i} \to e_i\} \rightsquigarrow \text{if app is}_{D_1} \, p \text{ then } p_1 \text{ else } \ldots \text{ else } p_n}
$$

The above translation yields the reflected definition for append (++) from (§ 2.5).

Semantic Preservation

$$\begin{array}{rrcl}
\textbf{\textit{Terms}} & p, t, b & ::= & \lambda^S \text{ if-free predicates from Figure 8} \\
\textbf{\textit{Functions}} & F & ::= & \lambda \overline{x}.\overline{\langle p \Rightarrow b \rangle} \\
\textbf{\textit{Definitional Environment}} & \Psi & ::= & \emptyset \mid f \mapsto F, \Psi \\
\textbf{\textit{Logical Environment}} & \Phi & ::= & \emptyset \mid p, \Phi
\end{array}$$

Fig. 9. Syntax of Predicates, Terms and Reflected Functions

We can prove that the translation preserves the semantics of the expressions. Informally, If $\Gamma \vdash e \rightsquigarrow p$, then for every substitution θ and every logical model σ that respects the environment Γ if $\theta \cdot e \hookrightarrow^\star v$ then $\sigma \models p = \lfloor v \rfloor$.

5.2 Algorithmic Type Checking

We make the type checking from Figure 7 algorithmic by checking subtyping via our novel, SMT-based *Proof by Logical Evaluation*(PLE). Next, we formalize how PLE makes checking algorithmic, and then describe it in detail in § 6.

Verification Conditions The implication or *verification condition* (VC) $\lfloor \Gamma \rfloor \Rightarrow p$ is *valid* only if the set of values described by Γ is subsumed by the set of values described by p. Γ is embedded into logic by conjoining the refinements of provably terminating binders [Vazou et al. 2014]:

$$\lfloor \Gamma \rfloor \doteq \bigcup_{x \in \Gamma} \lfloor \Gamma, x \rfloor \quad \text{where we embed each binder as} \quad \lfloor \Gamma, x \rfloor \doteq \begin{cases} \lfloor e \rfloor & \text{if } \Gamma(x) = \{x : B^\Downarrow \mid e\} \\ \text{True} & \text{otherwise.} \end{cases}$$

Validity Checking Given a reflection environment R, type environment Γ, and expression e, the procedure $\text{PLE}(\lfloor R \rfloor, \lfloor \Gamma \rfloor, \lfloor e \rfloor)$, returns *true* only when the expression e evaluates to True under the reflection and type environments R and Γ.

Subtyping via VC Validity Checking We make subtyping, and hence, typing decidable, by replacing the denotational base subtyping rule $\leq$-Base-λ^R with the conservative, algorithmic version $\leq$-Base-PBE that uses PLE to check the validity of the subtyping VC.

$$\frac{\text{PLE}(\lfloor R \rfloor, \lfloor \Gamma, v : \{v : B^\Downarrow \mid e\} \rfloor, \lfloor e' \rfloor)}{\Gamma; R \vdash_{PBE} \{v : B \mid e\} \preceq \{v : B \mid e'\}} \quad \leq\text{-Base-PBE}$$

This typing rule is sound as functions reflected in R always respect the typing environment Γ (by construction), and because PBE is sound (Theorem 6.2).

LEMMA 5.1. *If* $\Gamma; R \vdash_{PBE} \{v : B \mid e_1\} \preceq \{v : B \mid e_2\}$ *then* $\Gamma; R \vdash \{v : B \mid e_1\} \preceq \{v : B \mid e_2\}$.

Soundness of λ^S We write $\Gamma; R \vdash_S e : \tau$ for the judgments that can be derived by the algorithmic subtyping rule $\leq$-Base-λ^S (instead of $\leq$-Base-λ^R.) Lemma 5.1 implies the soundness of λ^S.

THEOREM 5.2 (SOUNDNESS OF λ^S). *If* $\Gamma; R \vdash_S e : \tau$ *then* $\Gamma; R \vdash e : \tau$.

6 COMPLETE VERIFICATION: PROOF BY LOGICAL EVALUATION

Next, we formalize our Proof By Logical Evaluation PLE algorithm and show that it is sound (§ 6.1), that it is complete with respect to equational proofs (§ 6.2), and that it terminates (§ 6.3).

Unfold	:	$(\Psi, \Phi) \to \Phi$
Unfold(Ψ, Φ)	=	$\Phi \cup \bigcup_{f(\bar{t}) \prec \Phi}$ Instantiate $(\Psi, \Phi, f, \bar{t})$
Instantiate $(\Psi, \Phi, f, \bar{t})$	=	$\left\{ (\lfloor f(\bar{x}) \rfloor = b_i) \left[\bar{t}/\bar{x} \right] \;\middle|\; (p_i \Rightarrow b_i) \in \bar{d}, \; \text{SmtValid}(\Phi, p_i \left[\bar{t}/\bar{x} \right]) \right\}$
where		
$\lambda \bar{x}.\langle \bar{d} \rangle$	=	$\Psi(f)$

PLE	:	$(\Psi, \Phi, p) \to \texttt{Bool}$
PLE(Ψ, Φ, p)	=	$\text{loop} \left(0, \Phi \cup \bigcup_{f(\bar{t}) \prec p} \text{Instantiate} \left(\Psi, \Phi, f, \bar{t} \right) \right)$
where		
loop (i, Φ_i)		
$\mid$ SmtValid(Φ_i, p)	=	*true*
$\mid \Phi_{i+1} \subseteq \Phi$	=	*false*
$\mid$ otherwise	=	loop $(i + 1, \Phi_{i+1})$
where		
Φ_{i+1}	=	$\Phi \cup \text{Unfold}(\Psi, \Phi_i)$

Fig. 10. Algorithm PLE: Proof by Logical Evaluation

6.1 Algorithm

Figure 9 describes the input environments for PBE. The logical environment Φ contains a set of hypotheses p, described in Figure 8. The definitional environment Ψ maps function symbols f to their definitions $\lambda \bar{x}.\langle p \Rightarrow b \rangle$, written as λ-abstractions over guarded bodies. Moreover, the body b and the guard p contain neither λ nor if. These restrictions do not impact expressiveness: λs can be named and reflected, and if-expressions can be pulled out into top-level guards using $\text{DeIf}(\cdot)$, found in Appendix ??. A definitional environment Ψ can be constructed from R as

$$\lfloor R \rfloor \;\dot{=}\; \{f \mapsto \lambda \bar{x}.\text{DeIf}(\lfloor e \rfloor) | (f \mapsto \lambda \bar{x}.e) \in R\}$$

Notation We write $f(\bar{t}) \prec \Phi$ if the λ^S term $(\text{app} \ldots (\text{app } f \ t_1) \ldots t_n)$ is a syntactic subterm of some $t' \in \Phi$. We abuse notation to write $f(\bar{t}) \prec t'$ for $f(\bar{t}) \prec \{t'\}$. We write $\text{SmtValid}(\Phi, p)$ for SMT validity of the implication $\Phi \Rightarrow p$.

Instantiation & Unfolding A term q is a (Ψ, Φ)-*instance* if there exists $f(\bar{t}) \prec \Phi$ such that:

- $\Psi(f) \equiv \lambda \bar{x}.\langle \overline{p_i \Rightarrow b_i} \rangle$,
- $\text{SmtValid}(\Phi, p_i \left[\bar{t}/\bar{x} \right])$,
- $q \equiv (f(\bar{x}) = b_i) \left[\bar{t}/\bar{x} \right]$.

A set of terms Q is a (Ψ, Φ)-*instance* if every $q \in Q$ is an (Ψ, Φ)-instance. The *unfolding* of Ψ, Φ is the (finite) set of all (Ψ, Φ)-instances. Procedure Unfold(Ψ, Φ) shown in Figure 10 computes and returns the conjunction of Φ and the unfolding of Ψ, Φ. The following properties relate (Ψ, Φ)-instances to the semantics of λ^R and SMT validity.

LEMMA 6.1. *For every $\Gamma \models R$, and $\theta \in (\!|\Gamma|\!)$,*

- ***Sat-Inst*** *If $\lfloor e \rfloor$ is a $(\lfloor R \rfloor, \lfloor \Gamma \rfloor)$-instance, then $\theta \cdot R[e] \hookrightarrow^\star \texttt{True}$.*
- ***SMT-Approx*** *If $\text{SmtValid}(\lfloor \Gamma \rfloor, \lfloor e \rfloor)$ then $\theta \cdot R[e] \hookrightarrow^\star \texttt{True}$.*
- ***SMT-Inst*** *If q is a $(\lfloor R \rfloor, \lfloor \Gamma \rfloor)$-instance and $\text{SmtValid}(\lfloor \Gamma \rfloor, q, \lfloor e \rfloor)$ then $\theta \cdot R[e] \hookrightarrow^\star \texttt{True}$.*

The Algorithm Figure 10 shows our proof search algorithm $\mathsf{PLE}(\Psi, \Phi, p)$ which takes as input a set of *reflected definitions* Ψ, an *hypothesis* Φ, and a *goal* p. The PBE procedure recursively *unfolds* function application terms by invoking Unfold until either the goal can be proved using the unfolded instances (in which case the search returns *true*) *or* no new instances are generated by the unfolding (in which case the search returns *false*).

Soundness First, we prove the soundness of PLE. Let $R[e]$ denote the evaluation of e under the reflection environment R, *i.e.* $\emptyset[e] \doteq e$ and $(R, f : e_f)[e] \doteq R[\text{let rec } f = e_f \text{ in } e]$.

> THEOREM 6.2 (**SOUNDNESS**). *If* $\mathsf{PLE}(\lfloor R \rfloor, \lfloor \Gamma \rfloor, \lfloor e \rfloor)$ *then* $\forall \theta \in \langle\!\langle \Gamma \rangle\!\rangle, \theta \cdot R[e] \hookrightarrow^{\star} \mathsf{True}$

We prove Theorem 6.2 using the Lemma 6.1 that relates instantiation, SMT validity, and the exact semantics. Intuitively, PLE is sound as it reasons about a finite set of instances by *conservatively* treating all function applications as *uninterpreted* [Nelson 1981].

6.2 Completeness

Next, we show that our proof search is *complete* with respect to equational reasoning. We define a notion of equational proof $\Psi, \Phi \vdash t \twoheadrightarrow t'$ and prove that if there exists such a proof, then $\mathsf{PLE}(\Psi, \Phi, t = t')$ is guaranteed to return *true*. To prove this theorem, we introduce the notion of *bounded unfolding* which corresponds to unfolding definitions n times. We will show that unfolding preserves congruences, and hence, that an equational proof exists iff the goal can be proved with *some* bounded unfolding. Thus, completeness follows by showing that the proof search procedure computes the limit (*i.e.* fixpoint) of the bounded unfolding. In § 6.3 we will show that the fixpoint is computable: there is an unfolding depth at which PLE reaches a fixpoint and hence terminates.

Bounded Unfolding For every Ψ, Φ and $0 \le n$, the *bounded unfolding of depth* n is defined by:

$$
\begin{aligned}
\mathsf{Unfold}^*(\Psi, \Phi, 0) &\doteq \Phi \\
\mathsf{Unfold}^*(\Psi, \Phi, n+1) &\doteq \Phi_n \cup \mathsf{Unfold}(\Psi, \Phi_n) \quad \text{where } \Phi_n = \mathsf{Unfold}^*(\Psi, \Phi, n)
\end{aligned}
$$

That is, the unfolding at depth n essentially performs Unfold upto n times. The bounded-unfoldings yield a monotonically non-decreasing sequence of formulas and that if two consecutive bounded unfoldings coincide, then all subsequent unfoldings are the same.

> LEMMA 6.3 (**MONOTONICITY**). $\forall 0 \le n. \mathsf{Unfold}^*(\Psi, \Phi, n) \subseteq \mathsf{Unfold}^*(\Psi, \Phi, n+1)$.

> LEMMA 6.4 (**FIXPOINT**). *Let* $\Phi_i \doteq \mathsf{Unfold}^*(\Psi, \Phi, i)$. *If* $\Phi_n = \Phi_{n+1}$ *then* $\forall n < m. \Phi_m = \Phi_n$.

Uncovering Next we prove that every function application term that is *uncovered* by unfolding to depth n is congruent to a term in the n-depth unfolding.

> LEMMA 6.5 (**UNCOVERING**). *Let* $\Phi_n \equiv \mathsf{Unfold}^*(\Psi, \Phi \cup \{v = t\}, n)$. *If* $\mathsf{SmtValid}(\Phi_n, v = t')$ *then for every* $f(\overline{t'}) \prec t'$ *there exists* $f(\overline{t}) \prec \Phi_n$ *such that* $\mathsf{SmtValid}(\Phi_n, t_i = t_i')$.

We prove the above lemma by induction on n where the inductive step uses the following property of congruence closure, which itself is proved by induction on the structure of t':

> LEMMA 6.6 (**CONGRUENCE**). *If* $\mathsf{SmtValid}(\Phi, v = t, v = t')$ *and* $v \notin \Phi, t, t'$ *then for every* $f(\overline{t'}) \prec t'$ *there exists* $f(\overline{t}) \prec \Phi, t$ *such that* $\mathsf{SmtValid}(\Phi, t_i = t_i')$.

Unfolding Preserves Equational Links Next, we use the uncovering Lemma 6.5 and congruence to show that *every instantiation* that is valid after n steps is subsumed by the $n + 1$ depth unfolding. That is, we show that every possible *link* in a possible equational chain can be proved equal to the source expression via bounded unfolding.

$$\overline{\Psi, \Phi \vdash t \twoheadrightarrow t} \quad \text{Eq-Refl}$$

$$\frac{\Psi, \Phi \vdash t \twoheadrightarrow t'' \quad \Phi' = \text{Unfold}(\Psi, \Phi \cup \{v = t''\}) \quad \text{SmtValid}(\Phi', v = t')}{\Psi, \Phi \vdash t \twoheadrightarrow t'} \quad \text{Eq-Trans}$$

$$\frac{\Psi, \Phi \vdash t_1 \twoheadrightarrow t_1' \quad \Psi, \Phi \vdash t_2 \twoheadrightarrow t_2' \quad \text{SmtValid}(\Phi, t_1' \bowtie t_2')}{\Psi, \Phi \vdash t \bowtie t'} \quad \text{Eq-Proof}$$

Fig. 11. Equational Proofs

LEMMA 6.7 (**LINK**). *If* $\text{SmtValid}(\text{Unfold}^*(\Psi, \Phi \cup \{v = t\}, n), v = t')$ *then* $\text{SmtValid}(\text{Unfold}^*(\Psi, \Phi \cup \{v = t\}, n + 1), \text{Unfold}(\Psi, \Phi \cup \{v = t'\}))$.

Equational Proof Figure 11 formalizes our rules for equational reasoning. Intuitively, there is an *equational proof* that $t_1 \bowtie t_2$ under Ψ, Φ written by the judgment $\Psi, \Phi \vdash t_1 \bowtie t_2$ if by some sequence of repeated function unfoldings, we can prove that t_1 and t_2 are respectively equal to t_1' and t_2' such that, $\text{SmtValid}(\Phi, t_1 \bowtie t_2)$ holds. Our notion of equational proofs adapts the idea of type level computation used in TT-based verifiers to the setting of SMT-based reasoning, via the directional unfolding judgment $\Psi, \Phi \vdash t \twoheadrightarrow t'$. In the SMT-realm, the explicit notion of a normal or canonical form is converted to the implicit notion of the equivalence classes of the SMT solver's congruence closure procedure (post-unfolding.)

Completeness of Bounded Unfolding Next, we use the fact that unfolding preserves equational links to show that bounded unfolding is *complete* for equational proofs. That is, we prove by induction on the structure of the equational proof that whenever there is an *equational proof* of $t = t'$, there exists some bounded unfolding that suffices to prove the equality.

LEMMA 6.8. *If* $\Psi, \Phi \vdash t \twoheadrightarrow t'$ *then* $\exists 0 \le n. \, \text{SmtValid}(\text{Unfold}^*(\Psi, \Phi \cup \{v = t\}, n), v = t')$.

PLE *is a Fixpoint of Bounded Unfolding* Next, we show that the proof search procedure PLE computes the least-fixpoint of the bounded unfolding and hence, returns *true* iff there exists *some* unfolding depth n at which the goal can be proved.

LEMMA 6.9 (**FIXPOINT**). $\text{PLE}(\Psi, \Phi, t = t')$ *iff* $\exists n. \, \text{SmtValid}(\text{Unfold}^*(\Psi, \Phi \cup \{v = t\}, n), v = t')$

The proof follows by observing that $\text{PLE}(\Psi, \Phi, t = t')$ computes the *least-fixpoint* of the sequence $\Phi_i \doteq \text{Unfold}^*(\Psi, \Phi', i)$ Specifically, we can prove by induction on i that at each invocation of loop (i, Φ_i) in Figure 10, Φ_i is equal to $\text{Unfold}^*(\Psi, \Phi \cup \{v = t\}, i)$, which then yields the result.

Completeness of PLE By combining Lemma 6.9 and Lemma 6.7 we can show that PLE is complete, *i.e.* if there is an equational proof that $t \bowtie t'$ under Ψ, Φ, then $\text{PLE}(\Psi, \Phi, t \bowtie t')$ returns *true*.

THEOREM 6.10 (**COMPLETENESS**). *If* $\Psi, \Phi \vdash t \bowtie t'$ *then* $\text{PLE}(\Psi, \Phi, t \bowtie t') = true$.

6.3 PLE Terminates

So far, we have shown that our proof search procedure PLE is both sound and complete. Both of these are easy to achieve simply by *enumerating* all possible instances and repeatedly querying the SMT solver. Such a monkeys-with-typewriters approach is rather impractical: it may never terminate. Fortunately, next, we show that in addition to being sound and complete with respect to equational proofs, if the hypotheses are transparent, then our proof search procedure always

terminates. Next, we describe transparency and explain intuitively why PLE terminates. We then develop the formalism needed to prove the termination theorem 6.17.

Transparency An environment Γ is *inconsistent* if SmtValid($\lfloor \Gamma \rfloor$, *false*). An environment Γ is *inhabited* if there exists some $\theta \in (\!|\Gamma|\!)$. We say Γ is *transparent* if it is either inhabited *or* inconsistent. As an example of a *non-transparent* Φ_0 consider the predicate lenA xs = 1 + lenB xs, where lenA and lenB are both identical definitions of the list length function. Clearly there is no θ that causes the above predicate to evaluate to *true*. At the same time, the SMT solver cannot (using the decidable, quantifier-free theories) prove a contradiction as that requires induction over xs. Thus, non-transparent environments are somewhat pathological, and in practice, we only invoke PLE on transparent environments. Either the environment is inconsistent, *e.g.* when doing a proof-by-contradiction, or *e.g.* when doing a proof-by-case-analysis we can easily find suitable concrete values via random [Claessen and Hughes 2000] or SMT-guided generation [Seidel et al. 2015].

Challenge: Connect Concrete and Logical Semantics As suggested by its name, the PLE algorithm aims to lift the notion of evaluation or computations into the level of the refinement logic. Thus, to prove termination, we must connect the two different notions of evaluation, the *concrete* (operational) semantics and the *logical* semantics being used by PLE. This connection is trickier than appears at first glance. In the concrete realm totality ensures that every reflected function f will terminate when run on any *individual* value v. However, in the logical realm, we are working with *infinite* sets of values, compactly represented via logical constraints. In other words, the logical realm can be viewed (informally) as an *abstract interpretation*, of the concrete semantics. We must carefully argue that despite the *approximation* introduced by the logical abstraction, the abstract interpretation will also terminate.

Solution: Universal Abstract Interpretation We make this argument in three parts. First, we formalize how PLE performs computation at the logical level via *logical steps* and *logical traces*. We show (Lemma 6.13) that the logical steps form a so-called *universal* (or *"must"*) abstraction of the concrete semantics [Clarke et al. 1992; Cousot and Cousot 1977]. Second, we show that if PLE diverges, it is because it creates a strictly increasing infinite chain, Unfold*($\Psi, \Phi, 0$) $\subset$ Unfold*($\Psi, \Phi, 1$) ... which corresponds to an *infinite logical trace*. Third, as the logical computation is universal abstraction we use inhabitation to connect the two realms, *i.e.* to show that an infinite logical trace corresponds to an infinite concrete trace. The impossibility of the latter must imply the impossibility of the former, and hence, PLE must terminate. Next, we formalize the above intuition to obtain Theorem 6.17.

Totality A function is *total* when its evaluation reduces to exactly one value. The totality of R can and is checked by refinement types [Vazou et al. 2014]. Hence, for brevity, in the sequel we will *implicitly assume* that R is total under Γ.

*Definition 6.11 (**Total**).* Let $b \equiv \lambda \overline{x}.\langle \overline{\lfloor p \rfloor \Rightarrow \lfloor e \rfloor} \rangle$. b is *total* under Γ and R if forall $\theta \in (\!|\Gamma|\!)$:

(1) If $\theta \cdot R[p_i] \hookrightarrow^\star$ True then $\theta \cdot R[e_i] \hookrightarrow^\star v$.

(2) If $\theta \cdot R[p_i] \hookrightarrow^\star$ True and $\theta \cdot \Psi[p_j] \hookrightarrow^\star$ True, then $i = j$.

(3) There exists an i so that $\theta \cdot R[p_i] \hookrightarrow^\star$ True.

R is *total* under Γ if every $b \in \lfloor R \rfloor$ is total under Γ and R.

Subterm Evaluation As the reflected functions are total, the Church-Rosser theorem implies that evaluation order is not important. To prove termination, we require an evaluation strategy, *e.g.* CBV, in which if a reflected function's guard is satisfied, then the evaluation of the corresponding function body requires evaluating *every subterm* inside the body. As DeIf($\cdot$) hoists if-expressions out of the body and into the top-level guards, the below fact follows from the properties of CBV:

LEMMA 6.12. *Let* $b \equiv \lambda \overline{x}.\langle \overline{\lfloor p \rfloor \Rightarrow \lfloor e \rfloor} \rangle$. *For every* Γ, R, *and* $\theta \in (\!|\Gamma|\!)$, *if* $\theta \cdot R[p_i] \hookrightarrow^\star$ True *and* $f(\overline{\lfloor e \rfloor}) \prec \lfloor e_i \rfloor$ *then* $\theta \cdot R[e_i] \hookrightarrow^\star C[f(\theta \cdot R[\overline{e}])]$.

Logical Step A pair $f(\overline{t}) \rightsquigarrow f'(\overline{t'})$ is a Ψ, Φ-*logical step* (abbrev. step) if

- $\Psi(f) \equiv \lambda \overline{x}.\langle \overline{p \Rightarrow b} \rangle$,
- SmtValid$(\Phi \wedge Q, p_i)$ for some (Ψ, Φ)-instance Q,
- $f'(\overline{t'}) \prec b_i \left[\overline{t}/\overline{x} \right]$

Steps and Reductions Next, using Lemmas 6.12, 6.1, and the definition of logical steps, we show that every logical step corresponds to a *sequence* of steps in the concrete semantics:

LEMMA 6.13 (**STEP-REDUCTIONS**). *If* $f(\overline{\lfloor e \rfloor}) \rightsquigarrow f'(\overline{\lfloor e' \rfloor})$ *is a logical step under* $\lfloor R \rfloor, \lfloor \Gamma \rfloor$ *and* $\theta \in (\!|\Gamma|\!)$, *then* $f(\overline{\theta \cdot R[e]}) \hookrightarrow^\star C[f(\theta \cdot R[\overline{e'}])]$ *for some context* C.

Steps and Values Next, we show that if $f(\overline{y}) \rightsquigarrow t'$ is a logical step under an Γ that is inhabited by θ then $f(\overline{y})$ reduces to a value under θ. The proof follows by observing that if Γ is inhabited by θ, and a particular step is possible, then the guard corresponding to that step must also be true under θ and hence, by totality, the function must reduce to a value under the given store.

LEMMA 6.14 (**STEP-VALUE**). *If* $\theta \in (\!|\Gamma|\!)$ *and* $f(\overline{y}) \rightsquigarrow t'$ *is a* $\lfloor R \rfloor, \lfloor \Gamma \rfloor$ *step then* $R[\theta \cdot f(\overline{y})] \hookrightarrow^\star v$.

Logical Trace A sequence $f_0(\overline{t_0}), f_1(\overline{t_1}), f_2(\overline{t_2}), \ldots$ is a Ψ, Φ-*logical trace* (abbrev. trace) if $\overline{t_0} \equiv \overline{x_0}$ for some variables $\overline{x_0}$, and $f_i(\overline{t_i}) \rightsquigarrow f_{i+1}(\overline{t_{i+1}})$ is a Ψ, Φ-step, for each i. Our termination proof hinges upon the following key result: inhabited environments only have *finite* logical traces. We prove this result by contradiction. Specifically, we show by Lemma 6.13 that an infinite $(\lfloor R \rfloor, \lfloor \Gamma \rfloor)$-trace combined with fact that Γ is inhabited yields *at least one infinite concrete trace*, which contradicts totality. Hence, all the $(\lfloor R \rfloor, \lfloor \Gamma \rfloor)$ logical traces must be finite.

THEOREM 6.15 (**FINITE-TRACE**). *If* Γ *is inhabited then every* $(\lfloor R \rfloor, \lfloor \Gamma \rfloor)$-*trace is finite*.

Ascending Chains and Traces If unfolding Ψ, Φ yields an infinite chain $\Phi_0 \subset \ldots \subset \Phi_n \ldots$, then Ψ, Φ has an infinite logical trace. We construct the trace by selecting, at level i, (*i.e.* in Φ_i), an application term $f_i(\overline{t_i})$ that was created by unfolding an application term at level $i - 1$ (*i.e.* in Φ_{i-1}).

LEMMA 6.16 (**ASCENDING CHAINS**). *Let* $\Phi_i \doteq$ Unfold$^*(\Psi, \Phi, i)$. *If there exists an (infinite) ascending chain* $\Phi_0 \subset \ldots \subset \Phi_n \ldots$ *then there exists an (infinite) logical trace* $f_0(\overline{t_0}), \ldots, f_n(\overline{t_n}), \ldots$.

Logical Evaluation Terminates Finally, we prove that the proof search procedure PLE terminates. If PLE loops forever, there must be an infinite strictly ascending chain of unfoldings Φ_i, and hence, by Lemma 6.16, an infinite logical trace, which, by Theorem 6.15, is impossible.

THEOREM 6.17 (**TERMINATION**). *If* R *is transparent then* PLE$(\lfloor R \rfloor, \lfloor \Gamma \rfloor, p)$ *terminates*.

7 EVALUATION

We have implemented refinement reflection in LIQUID HASKELL. In this section, we evaluate our approach by using LIQUID HASKELL to verify a variety of deep specifications of Haskell functions drawn from the literature and categorized in Figure 12, totalling about 4,000 lines of specifications and proofs. Verification is fast: from 1 to 7s per file (module), which is about twice as much time as GHC takes to compile the corresponding module. Overall there is about one line of type specification (theorems) for three lines of code (implementations and proof). Next, we detail each of the five classes of specifications, illustrate how they were verified using refinement reflection, and discuss

	CATEGORY		LOC
I.	**Arithmetic**		
	Fibonacci	§ 2	48
	Ackermann	[Tourlakis 2008]	280
II.	**Algebraic Data Types**		
	Fold Universal	[Mu et al. 2009]	105
III.	**Typeclass Laws**	Fig 13	
	Monoid	Peano, Maybe, List	189
	Functor	Maybe, List, Id, Reader	296
	Applicative	Maybe, List, Id, Reader	578
	Monad	Maybe, List, Id, Reader	435
IV.	**Functional Correctness**		
	SAT Solver	[Casinghino et al. 2014]	133
	Unification	[Sjöberg and Weirich 2015]	200
V.	**Deterministic Parallelism**		
	Conc. Sets	§ ??	1677
	n-body	§ ??	435
	Par. Reducers	§ ??	349
TOTAL			4155

Fig. 12. **Summary of Case Studies**

the strengths and weaknesses of our approach. *All* of these proofs require refinement reflection, *i.e.* are beyond the scope of shallow refinement typing.

Proof Strategies. Our proofs use three building blocks, that are seamlessly connected via refinement typing:

- ***Un/folding*** definitions of a function f at arguments e1...en, which due to refinement reflection, happens whenever the term f e1 ... en appears in a proof. For exposition, we render the function whose un/folding is relevant as **f**;
- ***Lemma Application*** which is carried out by using the "because" combinator ($\cdot$) to instantiate some fact at some inputs;
- ***SMT Reasoning*** in particular, *arithmetic*, *ordering* and *congruence closure* which kicks in automatically (and predictably!), allowing us to simplify proofs by not having to specify, *e.g.* which subterms to rewrite.

7.1 Arithmetic Properties

The first category of theorems pertains to the textbook Fibonacci and Ackermann functions. In § 2 we proved that fib is increasing. We prove a higher order theorem that lifts increasing functions to monotonic ones:

```
fMono :: f:(Nat → Int) → fUp:(z:Nat → {f z ≤ f (z+1)}) → x:Nat → y:{x < y} → {f
    x ≤ f y}
```

By instantiating the function argument f of fMono with fib, we proved monotonicity of fib.

Monoid	
Left Id.	$\text{mempty}\ x\ \diamond\ \equiv x$
Right Id.	$x\ \diamond\ \text{mempty} \equiv x$
Assoc	$(x\ \diamond\ y)\ \diamond\ z \equiv x\ \diamond\ (y\ \diamond\ z)$
Functor	
Id.	$\text{fmap id}\ xs \equiv \text{id}\ xs$
Distr.	$\text{fmap}\ (g \circ h)\ xs \equiv (\text{fmap}\ g\ \circ\ \text{fmap}\ h)\ xs$
Applicative	
Id.	$\text{pure id}\ \circledast\ v \equiv v$
Comp.	$\text{pure}\ (\circ)\ \circledast\ u\ \circledast\ v\ \circledast\ w \equiv u\ \circledast\ (v\ \circledast\ w)$
Hom.	$\text{pure}\ f\ \circledast\ \text{pure}\ x \equiv \text{pure}\ (f\ x)$
Inter.	$u\ \circledast\ \text{pure}\ y \equiv \text{pure}\ (\$\ y)\ \circledast\ u$
Monad	
Left Id.	$\text{return}\ a \ggg f \equiv f\ a$
Right Id.	$m \ggg \text{return} \equiv m$
Assoc	$(m \ggg f) \ggg g \equiv m \ggg (\lambda x \to f\ x \ggg g)$

Fig. 13. **Summary of Verified Typeclass Laws**

```
fibMono :: n:Nat → m:{n < m} → {fib n ≤ fib m}
fibMono = fMono fib fibUp
```

Using higher order functions like fMono, we mechanized the proofs of the Ackermann function properties from [Tourlakis 2008]. We proved 12 equational and arithmetic properties using various proof techniques, including instantiation of higher order theorems (like fMono), proof by construction, and natural number and generalized induction.

7.2 Algebraic Data Properties

Fold Universality as encoded in Adga [Mu et al. 2009]. The following encodes the universal property of foldr

```
foldr_univ :: f:(a → b → b)
           → h:(L a → b)
           → e:b
           → ys:L a
           → base:{h [] == e }
           → step:(x:a → xs:[a] → {h (x:xs) == f x (h xs)})
           → {h ys == foldr f e ys}
```

The Liquid Haskell proof term of foldr_univ is similar to the one in Agda. But, there are two major differences. First, specifications do not support quantifiers: universal quantification of x and xs in the step assumption is encoded as functional arguments. Second, unlike Agda, Liquid Haskell does not support optional arguments, thus to use foldr_univ the proof arguments base and step that were implicit arguments in Agda need to be explicitly provided. For example, the following code calls foldr_universal to prove foldr_fusion by explicitly applying proofs for the base and the step arguments

```
foldr_fusion :: h:(b → c)
```

```
                    → f:(a → b → b)
                    → g:(a → c → c)
                    → e:b
                    → z:[a]
                    → x:a → y:b
                    → fuse: {h (f x y) == g x (h y)})
                    → {(h.foldr f e) z == foldr g (h e) z}

  foldr_fusion h f g e ys fuse = foldr_univ g (h . foldr f e) (h e) ys
                                    (fuse_base h f e)
                                    (fuse_step h f e g fuse)
```

Where, for example, `fuse_base` is a function with type

```
  fuse_base :: h:(b → c) → f:(a → b → b) → e:b → {(h . foldr f e)
    [] == h e}
```

7.3 Typeclass Laws

We used LIQUID HASKELL to prove the Monoid, Functor, Applicative, and Monad Laws, summarized in Figure 13, for various user-defined instances summarized in Figure 12.

Monoid Laws A Monoid is a datatype equipped with an associative binary operator ◇ (mappend) and an *identity* element mempty. We prove that Peano (with a suitable add and Z), Maybe (with a suitable mappend and Nothing), and List (with append ++ and []) satisfy the monoid laws.

Functor Laws A type is a functor if it has a function fmap that satisfies the *identity* and *distribution* (or fusion) laws in Figure 13. For example, consider the proof of the map distribution law for the lists, also known as "map-fusion", which is the basis for important optimizations in GHC [Wiki 2008]. We reflect the definition of map for lists and use it to specify fusion and verify it by an inductive proof:

```
  map_fusion :: f:(b → c) → g:(a → b) → xs:[a] → {map (f . g) xs = (map f . map g)
    xs}
```

Monad Laws The monad laws, which relate the properties of the two operators ≫= and return (Figure 13), refer to λ-functions which are encoded in our logic as uninterpreted functions. For example, we can encode the associativity property as a refinement type alias:

```
  type AssocLaw m f g = { m >>= f >>= g = m >>= (\x → f x >>= g) }
```

and use it to prove that the list-bind is associative:

```
  assoc :: m:[a] → f:(a →[b]) → g:(b →[c]) → AssocLaw m f g
  assoc [] f g      = [] >>= f >>= g
                    ==== [] >>= g
                    ==== []
                    ==== [] >>= (\x → f x >>= g) ** QED
  assoc (x:xs) f g = (x:xs) >>= f  >>= g
                    ==== (f x ++ xs >>= f) >>= g
                    ==== (f x >>= g) ++ (xs >>= f >>= g) ∵ bind_append (f x) (xs >>= f)
    g
                    ==== (f x >>= g) ++ (xs >>= \y → f y >>= g) ∵ assoc xs f g
                    ==== (\y → f y >>= g) x ++ (xs >>= \y → f y >>= g) ∵ βeq f g x
```

```
==== (x:xs) >>= (\y → f y >>= g) ** QED
```

Where the `bind_append` fusion lemma states that:

```
bind_append
  :: xs:[a] → ys:[a] → f:(a → [b]) → {(xs++ys) >>= f = (xs >>= f)++(
     ys >>= f)}
```

Notice that the last step requires β-equivalence on anonymous functions, which we get by explicitly inserting the redex in the logic, via the following lemma with `trivial` proof

```
βeq :: f:_ → g:_ → x:_ → {f x >>= g = (\y → f y >>= g) x}
βeq _ _ _ = trivial
```

7.4 Functional Correctness

Next, we proved correctness of two programs from the literature: a unification algorithm and a SAT solver.

Unification We verified the unification of first order terms, as presented in [Sjöberg and Weirich 2015]. First, we define a predicate alias for when two terms `s` and `t` are equal under a substitution `su`:

```
eq_sub su s t = apply su s == apply su t
```

Now, we defined a Haskell function `unify s t` that can diverge, or return `Nothing`, or return a substitution `su` that makes the terms equal:

```
unify :: s:Term → t:Term → Maybe {su | eq_sub su s t}
```

For specification and verification, we only needed to reflect `apply` and not `unify`; thus, we only had to verify that the former terminates, and not the latter. We proved correctness by invoking separate helper lemmata. For example, to prove the post-condition when unifying a variable `TVar i` with a term `t` in which `i` *does not* appear, we apply a lemma `not_in`:

```
unify (TVar i) t2 | not (i ∈ freeVars t2) = Just (const [(i, t2)] ∵ not_in i t2)
```

i.e. if `i` is not free in `t`, the singleton substitution yields `t`:

```
not_in :: i:Int → t:{Term | not (i ∈ freeVars t)} → {eq_sub [(i,t)] (
     TVar i) t}
```

This example highlights the benefits of partial verification on a legacy programming language: potential diverging code (*e.g.* the function `unify`) coexists and invokes proof terms.

SAT Solver As another example, we implemented and verified the simple SAT solver used to illustrate and evaluate the features of the dependently typed language Zombie [Casinghino et al. 2014]. The solver takes as input a formula `f` and returns an assignment that *satisfies* `f` if one exists, as specified below.

```
solve :: f:Formula → Maybe {a:Assignment | sat a f}
```

The function `sat a f` returns `True` iff the assignment `a` satisfies the formula `f`. Verifying `solve` follows directly by reflecting `sat` into the refinement logic.

8 RELATED WORK

SMT-Based Verification SMT-solvers have been extensively used to automate program verification via Floyd-Hoare logics [Nelson 1981]. Our work is inspired by Dafny's Verified Calculations [Leino and Polikarpova 2016], and a framework for proving theorems in Dafny [Leino 2010], but differs in (1) our use of reflection instead of axiomatization and (2) our use of refinements to compose proofs. (3) our use of PLE to automate reasoning about user-defined functions. DAFNY (and the related F* [Swamy et al. 2016]) encode user-functions as axioms, and use a fixed fuel to instantiate functions upto some fixed unfolding depth [Amin et al. 2014]. While the fuel-based approach is incomplete, even for equational or calculational reasoning, it may, in practice be more expedient to quickly time out after a fixed, small number of instantiations rather than to perform an exhaustive proof search like PLE. Nevertheless, PLE demonstrates that it is possible to develop complete and practical algorithms for reasoning about user-defined functions in a decidable manner.

Dependent Types in Proof Assistants Reflection shows how deep specification and verification in the style of COQ [Bertot and Castéran 2004] and AGDA [Norell 2007] can be retrofitted into existing languages via refinement typing and PLE shows how type-level computation can be made compatible with SMT solvers' native theory reasoning yielding a powerful new way to automate proofs (§ 2.5). It would be interesting to investigate how the tactics and scriptable proof search of mature proof assistants like COQ and ISABELLE could be adapted to the refinement setting.

Dependent Types in Programming Integration of dependent types into Haskell has been a long standing goal that dates back to Cayenne [Augustsson 1998], a Haskell-like, fully dependent type language with undecidable type checking. In a recent line of work Eisenberg and Stolarek [2014] aim to allow fully dependent programming within Haskell, by making "type-level programming ... at least as expressive as term-level programming". Our approach differs significantly in that reflection and PLE allows SMT-aided verification, which drastically simplifies proofs over key theories like linear arithmetic and equality. Zombie [Casinghino et al. 2014; Sjöberg and Weirich 2015] investigates the design of a dependently typed language where SMT-style congruence closure is used to reason about the equality of terms. However, Zombie explicitly eschews type-level computation as the authors note that type level computation, or to be precise "equalities that follow from β-reduction" are "incompatible with congruence closure". The programmer must put in explicit `join` terms to indicate where computation should be triggered, and even here, equality checking is based on "fuel" and hence, incomplete. In contrast, reflection and PLE show how congruence closure and computation can be very profitably combined, freeing the programmer from having to provide explicit `join` operations.

Proving Equational Properties Several authors have proposed tools for proving (equational) properties of (functional) programs. Systems Sousa and Dillig [2016] and Asada et al. [2015] extend classical safety verification algorithms, respectively based on Floyd-Hoare logic and Refinement Types, to the setting of relational or k-safety properties that are assertions over k-traces of a program. Thus, these methods can automatically prove that certain functions are associative, commutative *etc.*. but are restricted to first-order properties and are not programmer-extensible. Zeno [Sonnex et al. 2012] generates proofs by term rewriting and Halo [Vytiniotis et al. 2013] uses an axiomatic encoding to verify contracts. Both the above are automatic, but unpredictable and not programmer-extensible, hence, have been limited to far simpler properties than the ones checked here. HERMIT [Farmer et al. 2015] proves equalities by rewriting the GHC core language, guided by user specified scripts. In contrast, our proofs are simply Haskell programs, we can use SMT solvers to automate reasoning, and, most importantly, we can connect the validity of proofs with the semantics of the programs.

9 CONCLUSIONS AND FUTURE WORK

We have shown how refinement reflection—namely reflecting the definitions of functions in their output refinements—can be used to convert a legacy programming language, like Haskell, into a theorem prover. Reflection ensures that (refinement) type checking stays decidable and predictable via careful design of the logic and proof combinators. Reflection enables programmers working with the highly tuned libraries, compilers and run-times of legacy languages to specify and verify arbitrary properties of their code simply by writing programs in the legacy language. Our evaluation shows that refinement reflection lets us prove deep specifications of a variety of implementations, for example, the determinism of fast parallel programming libraries, and also identifies two important avenues for research.

First, while our proofs are often elegant and readable, they can sometimes be cumbersome. For example, in the proof of associativity of the monadic bind operator for the Reader monad three out of eight (extensional) equalities required explanations, some nested under multiple λ-abstractions. Thus, it would be valuable to explore how to extend the notions of tactics, proof search, and automation to the setting of legacy languages. Similarly, while our approach to α- and β-equivalence is sound, we do not know if it is *complete*. We conjecture it is, due to the fact that our refinement terms are from the simply typed lambda calculus (STLC). Thus, it would be interesting to use the normalization of STLC to develop a sound and complete system for SMT-based type-level computation and use it to automate proofs predictably.

Second, while Haskell's separation of pure and effectful code undoubtedly makes it easier to implement refinement reflection, we believe that the technique, like refinement typing in general, is orthogonal to purity. In particular, by carefully mediating between the interaction of pure and impure code using methods like permissions, uniqueness *etc.* refinement type systems have been developed for legacy languages like C [Rondon et al. 2010], JavaScript [Chugh et al. 2012; Vekris et al. 2016], and Racket [Kent et al. 2016] and so it would be interesting to see how to extend refinement reflection to other legacy languages.

REFERENCES

N. Amin, K. R. M. L., and T. Rompf. 2014. Computing with an SMT Solver. In *TAP*.

Andrew Appel. 2016. *Verified Functional Algorithms*. https://www.cs.princeton.edu/~appel/vfa/Perm.html.

K. Asada, R. Sato, and N. Kobayashi. 2015. Verifying Relational Properties of Functional Programs by First-Order Refinement. In *PEPM*.

L. Augustsson. 1998. Cayenne - a Language with Dependent Types.. In *ICFP*.

C. Barrett, A. Stump, and C. Tinelli. 2010. The SMT-LIB Standard: Version 2.0.

J. Bengtson, K. Bhargavan, C. Fournet, A.D. Gordon, and S. Maffeis. 2008. Refinement Types for Secure Implementations. In *CSF*.

Y. Bertot and P. Castéran. 2004. *Coq'Art: The Calculus of Inductive Constructions*. Springer Verlag.

Richard S. Bird. 1989. Algebraic Identities for Program Calculation. *Comput. J.* 32, 2 (1989), 122–126. https://doi.org/10.1093/comjnl/32.2.122

C. Casinghino, V. Sjöberg, and S. Weirich. 2014. Combining proofs and programs in a dependently typed language. In *POPL*.

R. Chugh, D. Herman, and R. Jhala. 2012. Dependent Types for JavaScript. In *OOPLSA*.

Koen Claessen and John Hughes. 2000. QuickCheck: A Lightweight Tool for Random Testing of Haskell Programs. In *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming (ICFP '00)*. ACM, New York, NY, USA, 268–279. https://doi.org/10.1145/351240.351266

E. M. Clarke, O. Grumberg, and D.E. Long. 1992. Model checking and abstraction. In *POPL 92: Principles of Programming Languages*. ACM, 343–354.

J. Condit, M. Harren, Z. R. Anderson, D. Gay, and G. C. Necula. 2007. Dependent Types for Low-Level Programming. In *ESOP*.

R. L. Constable and S. F. Smith. 1987. Partial Objects In Constructive Type Theory. In *LICS*.

P. Cousot and R. Cousot. 1977. Abstract interpretation: a unified lattice model for the static analysis of programs. In *POPL 77*. ACM, 238–252.

E.W. Dijkstra. 1976. *A Discipline of Programming.* Prentice-Hall.

Edsger W. Dijkstra. 1975. Guarded Commands, Nondeterminacy and Formal Derivation of Programs. *Commun. ACM* 18, 8 (Aug. 1975), 453–457. https://doi.org/10.1145/360933.360975

R. A. Eisenberg and J. Stolarek. 2014. Promoting functions to type families in Haskell. In *Haskell*.

A. Farmer, N. Sculthorpe, and A. Gill. 2015. Reasoning with the HERMIT: Tool Support for Equational Reasoning on GHC Core Programs *(Haskell)*.

G. Gentzen. 1935. Inverstigations into Logical Deduction. (1935).

W. A. Howard. 1980. The formulae-as-types notion of construction. (1980). http://lecomte.al.free.fr/ressources/PARIS8_LSL/Howard80.pdf

A. M. Kent, D. Kempe, and S. Tobin-Hochstadt. 2016. Occurrence typing modulo theories. In *PLDI*.

K.W. Knowles and C. Flanagan. 2010. Hybrid type checking. *TOPLAS* (2010).

K. R. M. Leino. 2010. Dafny: An Automatic Program Verifier for Functional Correctness *(LPAR)*.

K. R. M. Leino and C. Pit-Claudel. 2016. Trigger selection strategies to stabilize program verifiers. In *CAV*.

K. R. M. Leino and N. Polikarpova. 2016. Verified Calculations. In *VSTTE*.

Rustan Leino. 2016. Dafny. (2016). https://github.com/Microsoft/dafny/blob/master/Test/dafny0/Fuel.dfy.

S. C. Mu, H. S. Ko, and P. Jansson. 2009. Algebra of Programming in Agda: Dependent Types for Relational Program Derivation. *J. Funct. Program.* (2009).

G. Nelson. 1981. *Techniques for program verification.* Technical Report CSL81-10. Xerox Palo Alto Research Center.

U. Norell. 2007. *Towards a practical programming language based on dependent type theory.* Ph.D. Dissertation. Chalmers.

X. Ou, G. Tan, Y. Mandelbaum, and D. Walker. 2004. Dynamic Typing with Dependent Types. In *IFIP TCS*.

J. C. Reynolds. 1972. Definitional interpreters for higher-order programming languages. In *25th ACM National Conference*.

P. Rondon, M. Kawaguchi, and R. Jhala. 2008. Liquid Types. In *PLDI*.

P. Rondon, M. Kawaguchi, and R. Jhala. 2010. Low-Level Liquid Types. In *POPL*.

J. Rushby, S. Owre, and N. Shankar. 1998. Subtypes for Specifications: Predicate Subtyping in PVS. *IEEE TSE* (1998).

Eric L. Seidel, Niki Vazou, and Ranjit Jhala. 2015. Type Targeted Testing. In *Proceedings of the 24th European Symposium on Programming on Programming Languages and Systems - Volume 9032*. Springer-Verlag New York, Inc., New York, NY, USA, 812–836. https://doi.org/10.1007/978-3-662-46669-8_33

V. Sjöberg and S. Weirich. 2015. Programming Up to Congruence. *POPL* (2015).

W. Sonnex, S. Drossopoulou, and S. Eisenbach. 2012. Zeno: An Automated Prover for Properties of Recursive Data Structures. In *TACAS*.

M. Sousa and I. Dillig. 2016. Cartesian hoare logic for verifying k-safety properties. In *PLDI*.

Non-Anonymous Supplementary-Material. 2017. Supplementary Material. (2017).

N. Swamy, C. Hriţcu, C. Keller, A. Rastogi, A. Delignat-Lavaud, S. Forest, K. Bhargavan, C. Fournet, P. Y. Strub, M. Kohlweiss, J. K. Zinzindohoue, and S. Zanella-Béguelin. 2016. Dependent Types and Multi-Monadic Effects in F*. In *POPL*.

G. Tourlakis. 2008. AckermannâĂŹs Function. (2008). http://www.cs.yorku.ca/~gt/papers/Ackermann-function.pdf.

N. Vazou, P. Rondon, and R. Jhala. 2013. Abstract Refinement Types. In *ESOP*.

N. Vazou, E. L. Seidel, R. Jhala, D. Vytiniotis, and S. Peyton-Jones. 2014. Refinement Types for Haskell. In *ICFP*.

P. Vekris, B. Cosman, and R. Jhala. 2016. Refinement types for TypeScript. In *PLDI*.

D. Vytiniotis, S.L. Peyton-Jones, K. Claessen, and D. Rosén. 2013. HALO: haskell to logic through denotational semantics. In *POPL*.

Philip Wadler. 1987. A Critique of Abelson and Sussman or Why Calculating is Better Than Scheming. *SIGPLAN Not.* 22, 3 (March 1987), 83–94. https://doi.org/10.1145/24697.24706

GHC Wiki. 2008. GHC Optimisations. (2008). https://wiki.haskell.org/GHC_optimisations.

H. Xi and F. Pfenning. 1998. Eliminating Array Bound Checking Through Dependent Types.. In *PLDI*.

A PROOF OF MAPFUSION WITHOUT PLE

```
  map_fusion f g []
  =  (map f . map g) []
  ====  map f (map g [])
  ====  map f []
  ====  map (f . g) []
  **  QED
map_fusion f g (C x xs)
  =  map (f . g) (x : xs)
```

```
==== ((f . g) x) : (map (f . g) xs)
==== ((f . g) x) : (((map f) . (map g)) xs)
   ? map_fusion f g xs
==== ((f . g) x) : (map f (map g xs))
==== (f (g x)) : (map f (map g xs))
==== map f ((g x) : (map g xs))
==== map f ((g x) : (map g xs))
==== map f (map g (x : xs))
==== map f ((map g) (x : xs))
==== ((map f) . (map g)) (x : xs)
** QED
```

B PROOFS

Proofs for § 6

PROOF. (Of lemma 6.5) We prove the result by induction on n.

Case $n = 0$: Immediate as $t \equiv t'$.

Case $n = k + 1$: Consider any t' such that $\mathsf{SmtValid}(\Phi_{k+1}, v = t')$. By definition $\Phi_{k+1} = \mathsf{Unfold}(\Psi, \Phi_k)$, hence $\mathsf{SmtValid}(\mathsf{Unfold}(\Psi, \Phi_k), v = t, v = t')$. Consider any $f(\overline{t'}) \prec t'$; Lemma 6.6 completes the proof. □

PROOF. (Of lemma 6.6) Proof by induction on the structure of t': Case $t' \equiv x$: There are no subterms, hence immediate. Case $t' \equiv c$: There are no subterms, hence immediate. Case $t' \equiv f(\overline{t'})$: Consider the *last link* in Φ connecting the equivalence class of v (and t) to t'. Suppose the last link is a *congruence link* of the form $t = t'$ where $t \equiv f(\overline{t})$ and $\mathsf{SmtValid}(\Phi, t = t')$. Then $f(\overline{t}) \prec \Phi, t$ and we are done. Suppose instead, the last link is an *equality link* in Φ of the form $z = f(\overline{t'})$. In this case, $f(\overline{t'}) \prec \Phi$ and again, we are done. □

PROOF. (Of lemma 6.7)

Let $G_k \doteq \mathsf{Unfold}^*(\Psi, \Phi, k)$. Let us assume that

$$\mathsf{SmtValid}(\mathsf{Unfold}^*(\Psi, \Phi, v = t, n), v = t') \tag{1}$$

Consider any instance

$$q \equiv f(\overline{t'}) = b_i \left[\overline{t'}/\overline{x} \right] \text{ in } \mathsf{Unfold}(\Psi, \Phi \wedge v = t') \tag{2}$$

By the definition of Unfold, we have

$$f(\overline{t'}) \prec \Phi \wedge v = t' \text{ such that } \mathsf{SmtValid}(\Phi, p_i \left[\overline{t'}/\overline{x} \right]) \tag{3}$$

By (1) and Lemma 6.5 there exists $f(\overline{t}) \prec \Phi_n$ such that

$$\mathsf{SmtValid}(\Phi_n, t = t')$$

As $\Phi \subseteq \Phi_n$ and (3), by congruence

$$\mathsf{SmtValid}(\Phi_n, p_i \left[\overline{t}/\overline{x} \right])$$

Hence, the instance

$$f(\overline{t}) = b_i \left[\overline{t}/\overline{x} \right] \text{ is in} \Phi_{n+1}$$

That is

$$\mathsf{SmtValid}(\Phi_{n+1}, t = t' \wedge f(\overline{t}) = b_i [t/x])$$

And so by congruence closure

$$\text{SmtValid}(\Phi_{n+1}, q)$$

As the above holds for every instance, we have

$$\text{SmtValid}(\Phi_{n+1}, \text{Unfold}(\Psi, \ \Phi \wedge v = t'))$$

□

PROOF. (Of lemma 6.8) The proof follows by induction on the structure of $\Psi, \Phi \vdash t \twoheadrightarrow t'$.

Base Case EQ-REFL: Follows immediately as $t \equiv t'$.

Inductive Case EQ-TRANS

In this case, there exists t'' such that

$$\Psi, \Phi \vdash t \twoheadrightarrow t'' \tag{4}$$

$$\text{SmtValid}(\text{Unfold}(\Psi, \ \Phi \wedge v = t''), v = t') \tag{5}$$

By the induction hypothesis (4) implies there exists $0 \le n$ such that

$$\text{SmtValid}(\text{Unfold}^*(\Psi, \Phi \wedge v = t, n), v = t'')$$

By Lemma 6.7 we have

$$\text{SmtValid}(\text{Unfold}^*(\Psi, \Phi, v = t, n + 1), \text{Unfold}(\Psi, \ \Phi \wedge v = t''))$$

Thus, by (5) and modus ponens we get

$$\text{SmtValid}(\text{Unfold}^*(\Psi, \Phi, v = t, n + 1), v = t')$$

□

PROOF. (Of Lemma 6.9) Let $\Phi' = \Phi \wedge v = t$.

Case $\Rightarrow$: Assume that $\text{PLE}(\Psi, \Phi, t = t')$. That is, at some iteration i we have $\text{SmtValid}(\Phi_i, v = t')$, *i.e.* by (6.2) we have $\text{SmtValid}(\text{Unfold}^*(\Psi, \Phi', i), v = t')$.

Case $\Leftarrow$: Pick the smallest n such that $\text{SmtValid}(\text{Unfold}^*(\Psi, \Phi', n), v = t')$. Using Lemmas 6.3 and 6.4 we can then show that forall $0 \le k < n$, we have

$$\text{Unfold}^*(\Psi, \Phi', k) \subset \text{Unfold}^*(\Psi, \Phi', k + 1)$$

and

$$\text{Unfold}^*(\Psi, \Phi', k) \nvdash v = t'$$

Hence, after n iterations of the recursive loop, $\text{PLE}(\Psi, \Phi, t = t')$, returns *true*. □

PROOF. (Of Lemma 6.14)

$$\text{Assume that} \quad \theta \in (\![\Gamma]\!) \tag{6}$$

$$\text{Let} \quad \theta^* \doteq \theta[\overline{x} \mapsto \theta \cdot \overline{y}] \tag{7}$$

$$\Psi(f) \doteq \lambda \overline{x}.\overline{\langle p \Rightarrow b \rangle} \tag{8}$$

As $f(\overline{y}) \rightsquigarrow t'$ is a $\lfloor R \rfloor \Gamma$ step, for some i, $\lfloor R \rfloor$-instance Q we have

$$\text{SmtValid}(\lfloor \Gamma \rfloor \wedge Q, p_i [\overline{y}/\overline{x}])$$

$$\text{Hence, by (6) and Lemma 6.1} \quad \theta \cdot R[p_i [\overline{y}/\overline{x}]] \hookrightarrow^\star \text{True} \qquad\qquad (9)$$

$$\text{As} \quad \theta^* \cdot p_i \equiv \theta \cdot p_i [\overline{y}/\overline{x}] \qquad\qquad (10)$$

$$\text{The fact (9) yields} \quad \theta^* \cdot R[p_i] \hookrightarrow^\star \text{True}$$

$$\text{By the Totality Assumption 6.11} \quad R[\theta^* \cdot f(\overline{x})] \hookrightarrow^\star v$$

$$\text{That is} \quad R[\theta \cdot f(\overline{y})] \hookrightarrow^\star v \qquad\qquad (11)$$

$$(12)$$

$$\square$$

Divergence A closed term t *diverges under R* if there is no v such that $R[t] \hookrightarrow^\star v$.

LEMMA B.1 (**DIVERGENCE**). *If $\forall 0 \leq i$ we have $R[t_i] \hookrightarrow^\star C[t_{i+1}]$ then t_0 diverges under Ψ.*

PROOF. (Of Theorem 6.15)

$$\text{Assume that} \quad \theta \in (\!|\Gamma|\!) \qquad\qquad (13)$$

$$\text{and assume an infinite } \lfloor R \rfloor, \lfloor \Gamma \rfloor \text{ trace:} \quad f_0(\overline{t_0}), f_1(\overline{t_1}), \dots \qquad\qquad (14)$$

$$\text{Where additionally} \quad \overline{t_0} \equiv \overline{x_0} \qquad\qquad (15)$$

$$\text{Define} \quad t_i^* \equiv \theta \cdot f_i(\overline{t_i}) \qquad\qquad (16)$$

$$\text{By Lemma 6.13, for every } i \in \mathbb{N} \quad R[t_i^*] \hookrightarrow^\star C_i[t_{i+1}^*]$$

$$\text{Hence, by Lemma B.1} \quad t_0^* \text{ diverges under } \Psi$$

$$\textit{i.e., by (15, 16)} \quad \theta \cdot f_0(\overline{x_0}) \text{ diverges under } \Psi \qquad\qquad (17)$$

$$\text{But by (13) and Lemma 6.14} \quad R[\theta \cdot f_0(\overline{x_0})] \hookrightarrow^\star v \quad \text{contradicting (17)}$$

Hence, assumption (14) cannot hold, *i.e.* all all the Ψ, Φ symbolic traces must be finite. $\square$

PROOF. (Of Theorem 6.17) As Φ is transparent, there are two cases.

Case: Γ *is inconsistent.*

$$\text{By definition of inconsistency} \quad \text{SmtValid}(\lfloor \Gamma \rfloor, \textit{false})$$

$$\text{Hence} \quad \text{SmtValid}(\lfloor \Gamma \rfloor, p)$$

$$\text{That is} \quad \text{PLE}(\lfloor R \rfloor, \lfloor \Gamma \rfloor, p) \text{ terminates immediately.}$$

Case: Γ *is inhabited.*

$$\text{That is, exists } \theta \text{ s.t.} \quad \theta \in (\!|\Gamma|\!) \qquad\qquad (18)$$

Suppose that $\text{PLE}(\lfloor R \rfloor, \lfloor \Gamma \rfloor, p)$ does not terminate.

$$\text{That is, there is an infinitely increasing chain:} \quad \Phi_0 \subset \dots \subset \Phi_n \dots \qquad\qquad (19)$$

$$\text{By Lemma 6.16} \quad \lfloor R \rfloor, \lfloor \Gamma \rfloor \text{ has an infinite trace}$$

Which, by (18) contradicts Theorem 6.15. Thus, (19) is impossible, *i.e.* $\text{PLE}(\lfloor R \rfloor, \lfloor \Gamma \rfloor, p)$ terminates.

$$\square$$

C PROOF OF SECTION ??

Lemma C.1 (Validity). *If there exists $e \in (\!|\phi|\!)$ then ϕ is valid.*

Proof. We prove the lemma by case analysis in the shape of ϕ.

- $\phi \equiv \{p\}$. Since the set $(\!|\{p\}|\!) = \{e | p \hookrightarrow^\star \text{True}\}$ is not empty, then $p \hookrightarrow^\star \text{True}$.
- $\phi \equiv \phi_1 \to \phi_2$. By assumption, there exists an expressions f so that $\forall e_x \in (\!|\phi_1|\!), f\ e_x \in (\!|\phi_2|\!)$. So, if there exists an expression $e_1 \in (\!|\phi_1|\!)$ that makes ϕ_1 valid then $f\ e_1$ makes ϕ_2 valid.
- $\phi \equiv \phi \to \text{False}$. By assumption, there exists an expressions f so that $\forall e_x \in (\!|\phi|\!), f\ e_x \in (\!|\{\text{False}\}|\!)$. So, if there exists an expression $e_1 \in (\!|\phi|\!)$ that makes ϕ valid then $f\ e_1$ makes $\{\text{False}\}$ valid, which is impossible, thus ϕ cannot be valid.
- $\phi \equiv \text{PAnd } \phi_1\ \phi_2$. If there exists a total expression $e \in (\!|\phi|\!)$ then e evaluates to PAnd $e_1\ e_2$.
- $\phi \equiv \text{POr } \phi_1\ \phi_2$. If there exists a total expression $e \in (\!|\phi|\!)$ then e evaluates to either POrL e' or POrR e'.
- $\phi \equiv x : \tau \to \phi$. By assumption, there exists an expressions f so that $\forall e_x \in (\!|\tau|\!), f\ e_x \in (\!|\phi[x/e_x]|\!)$. So, if there exists an expression $e_1 \in (\!|\tau|\!)$ then $f\ e_1$ makes $\phi[x/e_1]$ valid.
- $\phi \equiv (x :: \tau, \phi)$. By assumption, there exists an expressions p that evaluates to a pair (e_x, e_y) so that $e_x \in (\!|\tau|\!)$ and $e_y \in (\!|\phi[x/e_x]|\!)$.

□

Theorem C.2 (Validity). *If $\emptyset; \emptyset \vdash e : \phi$ then ϕ is valid.*

Proof. By direct implication of Lemma C.1 and soundness of λ^R (Theorem 4.1. □